



From Shader Code to a **Teraflop**: How Shader Cores Work

Kayvon Fatahalian
Stanford University

SIGGRAPH2008

Beyond Programmable Shading: Fundamentals



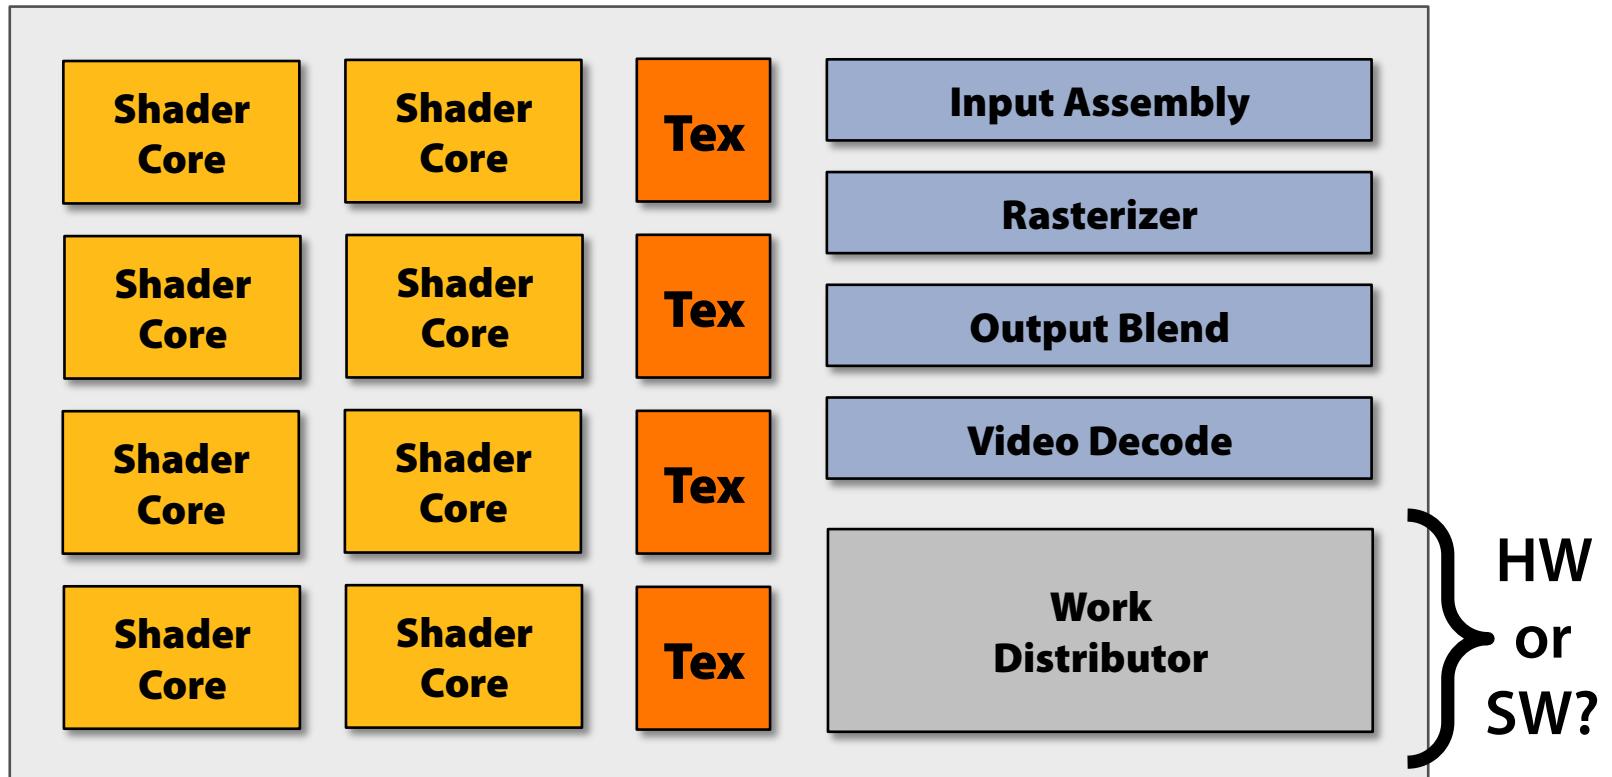
This talk

- Three key concepts behind how modern architectures run “shader” code
- Knowing these concepts will help you:
 1. Understand space of GPU shader core (and throughput CPU processing core) designs
 2. Optimize shaders/compute kernels
 3. Establish intuition: what workloads might benefit from the design of these architectures?



SIGGRAPH2008

What's in a GPU?



Heterogeneous chip multi-processor (highly tuned for graphics)

A diffuse reflectance shader



```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv)  
{  
    float3 kd;  
    kd = myTex.Sample(mySamp, uv);  
    kd *= clamp( dot(lightDir, norm), 0.0, 1.0);  
    return float4(kd, 1.0);  
}
```

Independent, but no explicit parallelism



SIGGRAPH2008

Compile shader

1 unshaded fragment input record



```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv)  
{  
    float3 kd;  
    kd = myTex.Sample(mySamp, uv);  
    kd *= clamp ( dot(lightDir, norm), 0.0,  
    1.0 );  
    return float4(kd, 1.0);  
}
```



```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul r3, v0, cb0[0]  
madd r3, v1, cb0[1], r3  
madd r3, v2, cb0[2], r3  
clmp r3, r3, 1(0.0), 1(1.0)  
mul o0, r0, r3  
mul o1, r1, r3  
mul o2, r2, r3  
mov o3, 1(1.0)
```



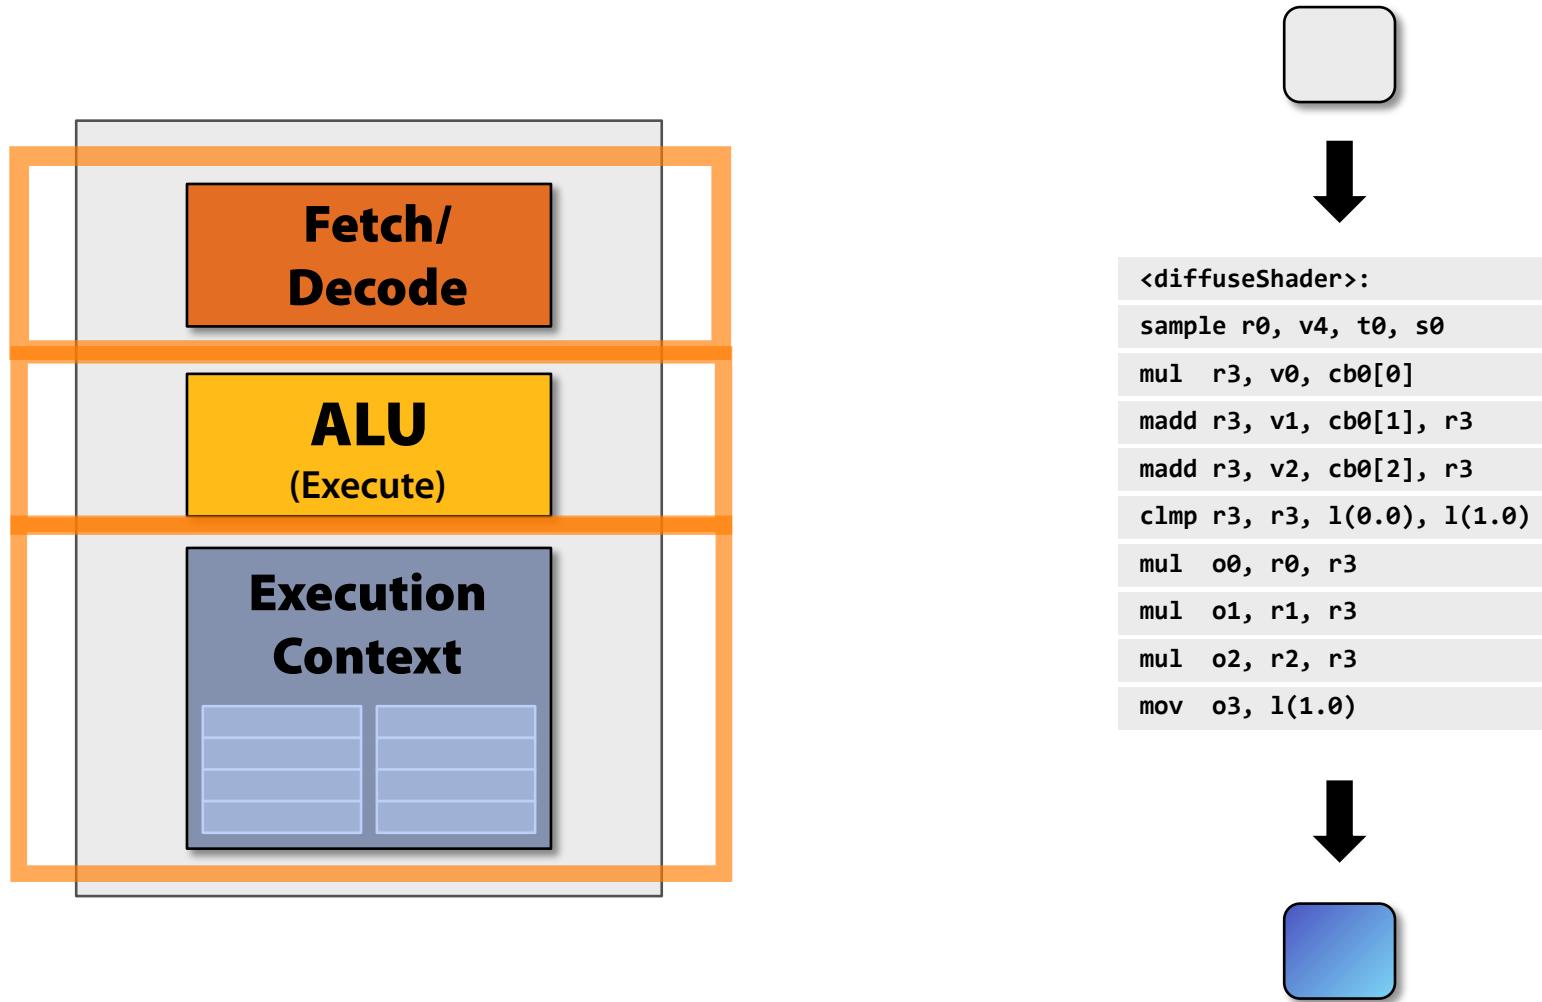
1 shaded fragment output record





SIGGRAPH2008

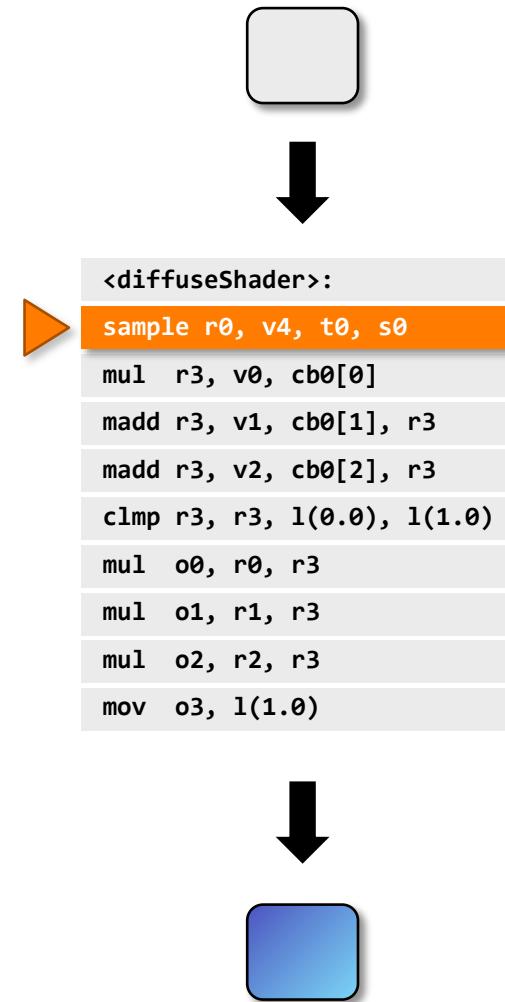
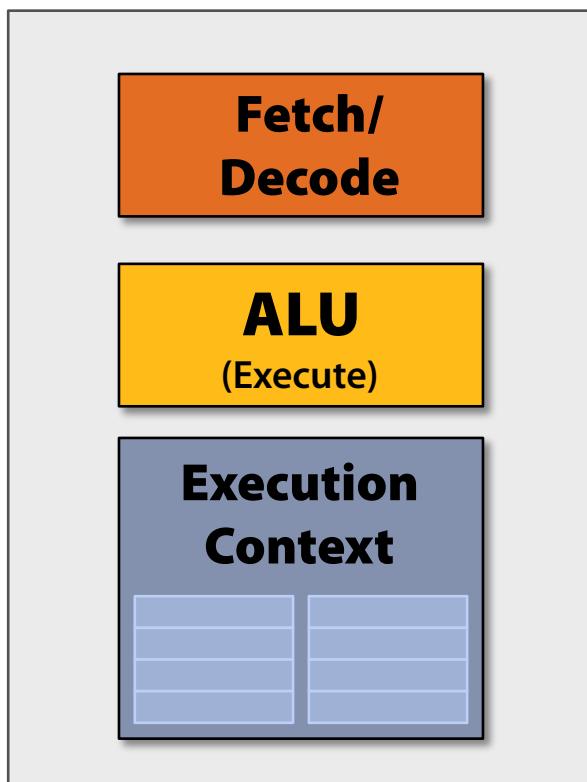
Execute shader





SIGGRAPH2008

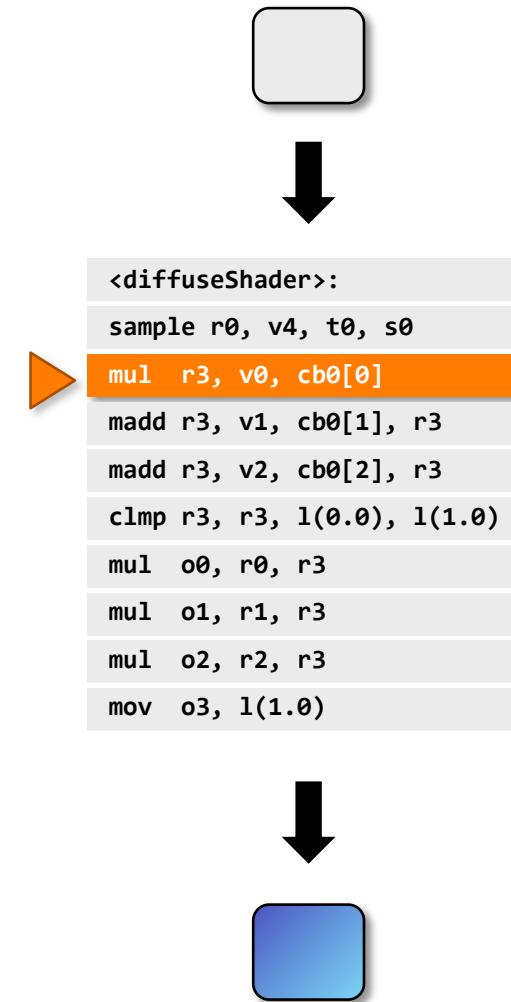
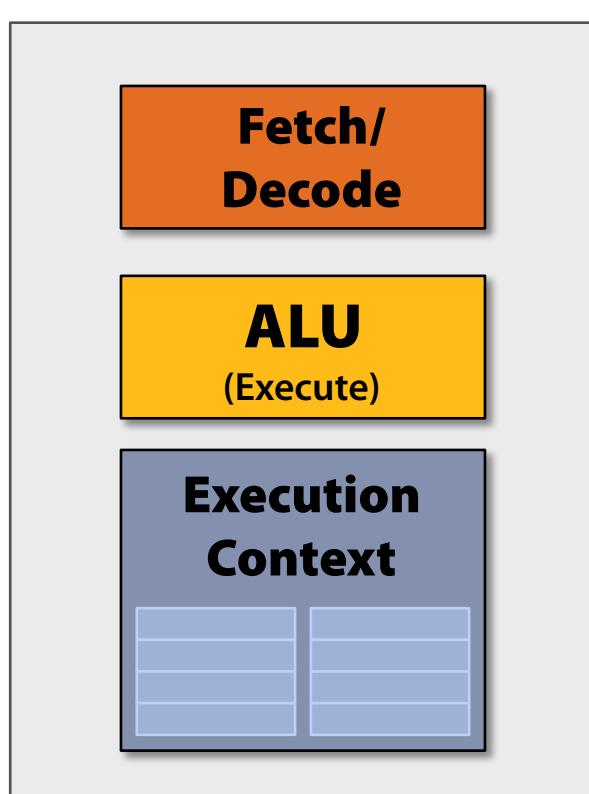
Execute shader





SIGGRAPH2008

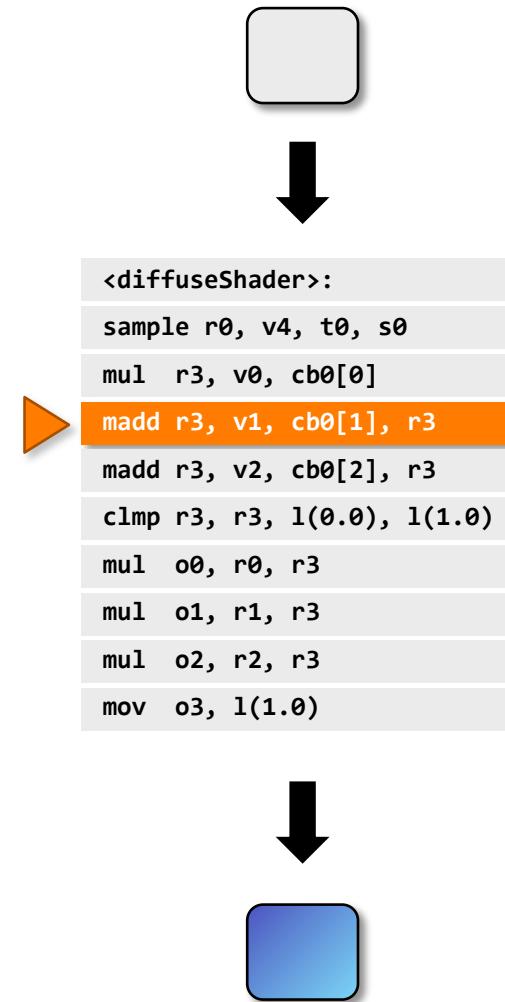
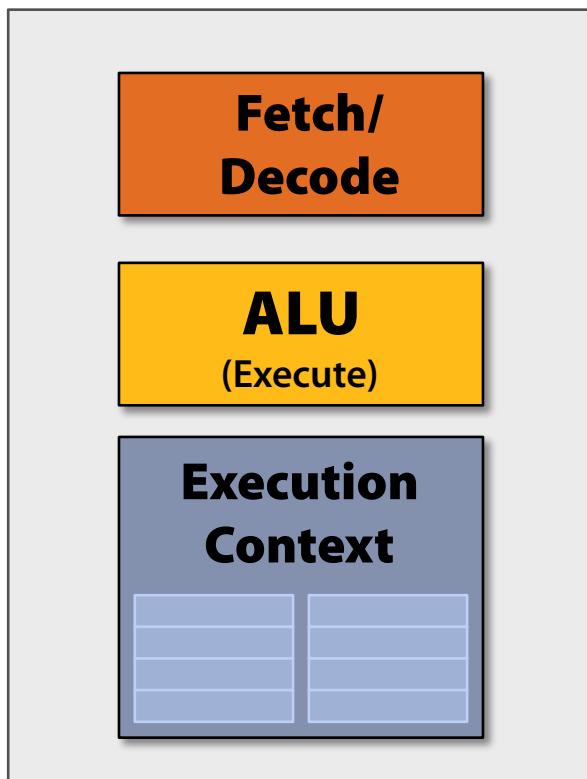
Execute shader





SIGGRAPH2008

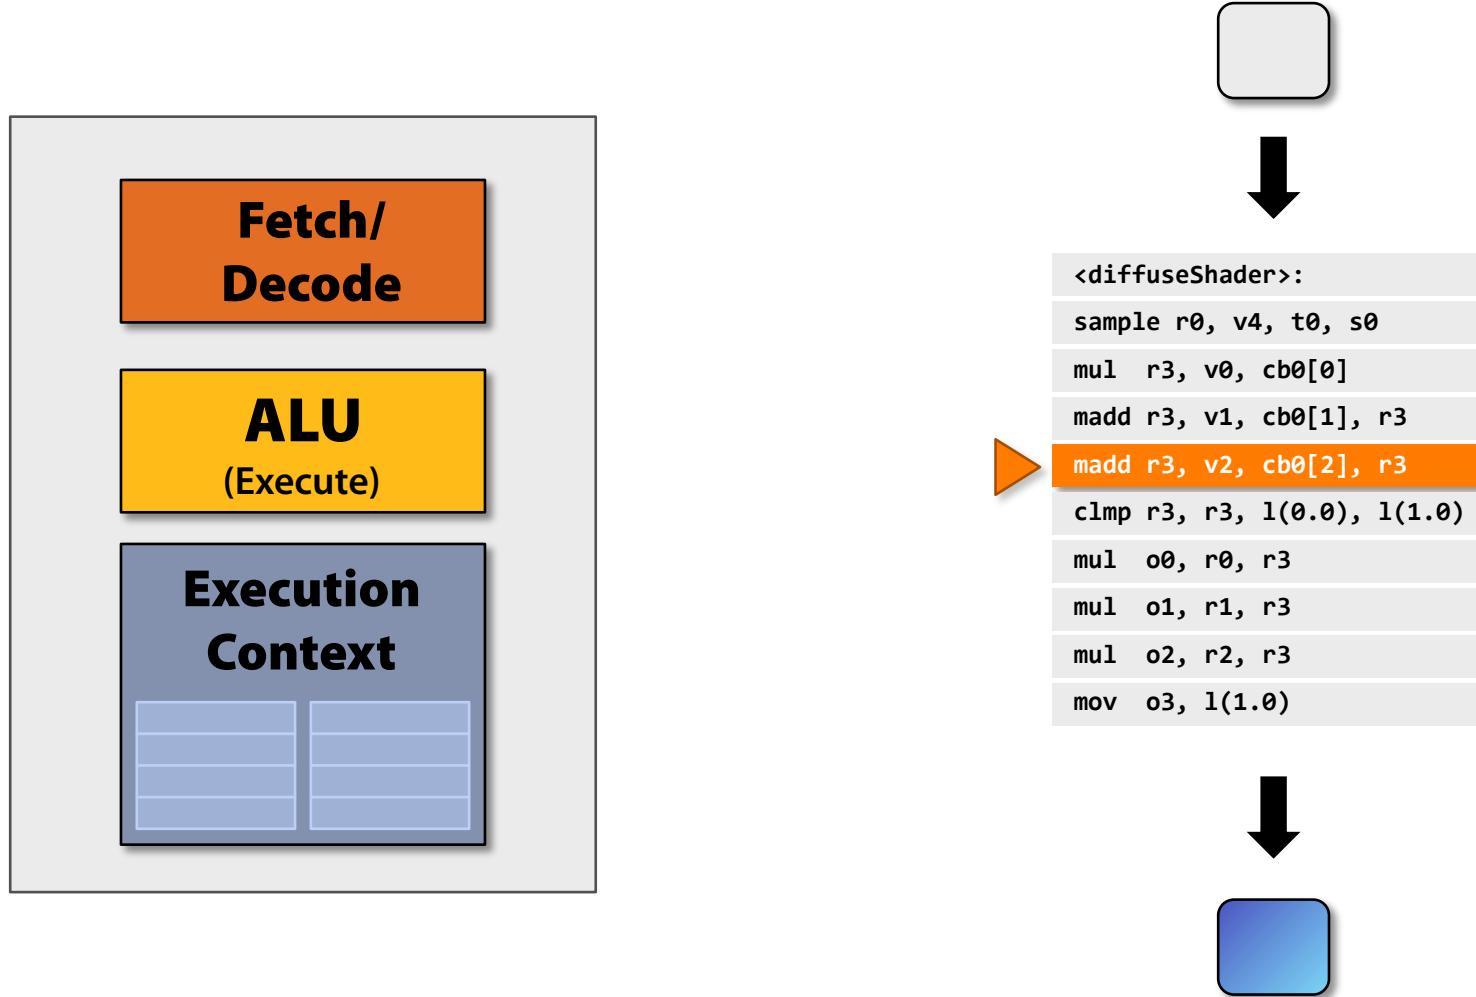
Execute shader





SIGGRAPH2008

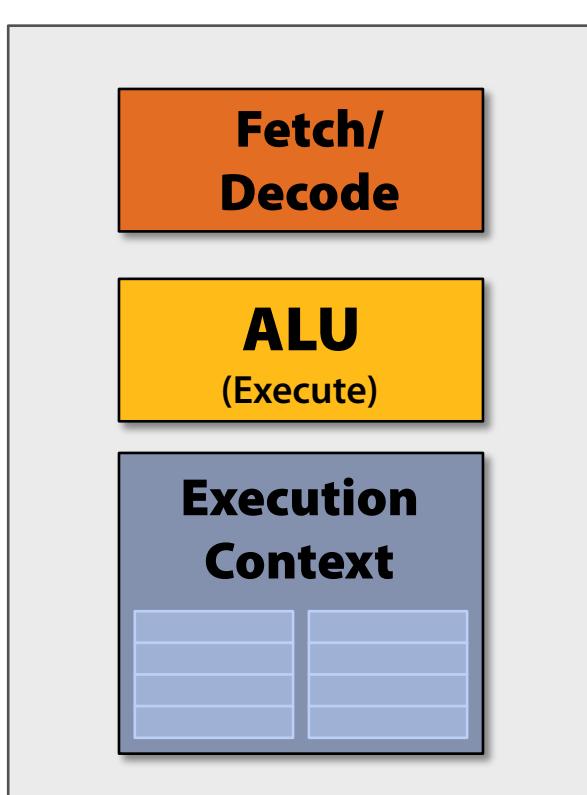
Execute shader





SIGGRAPH2008

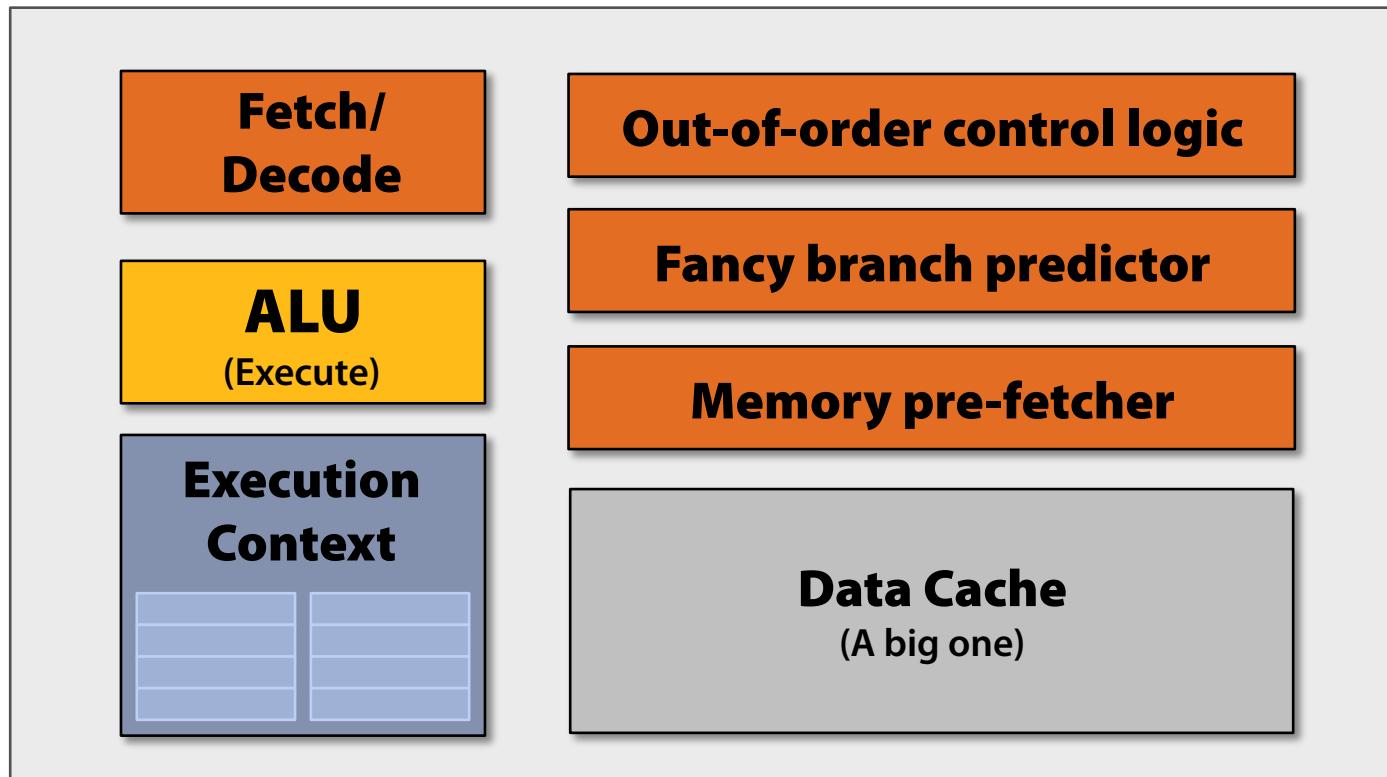
Execute shader



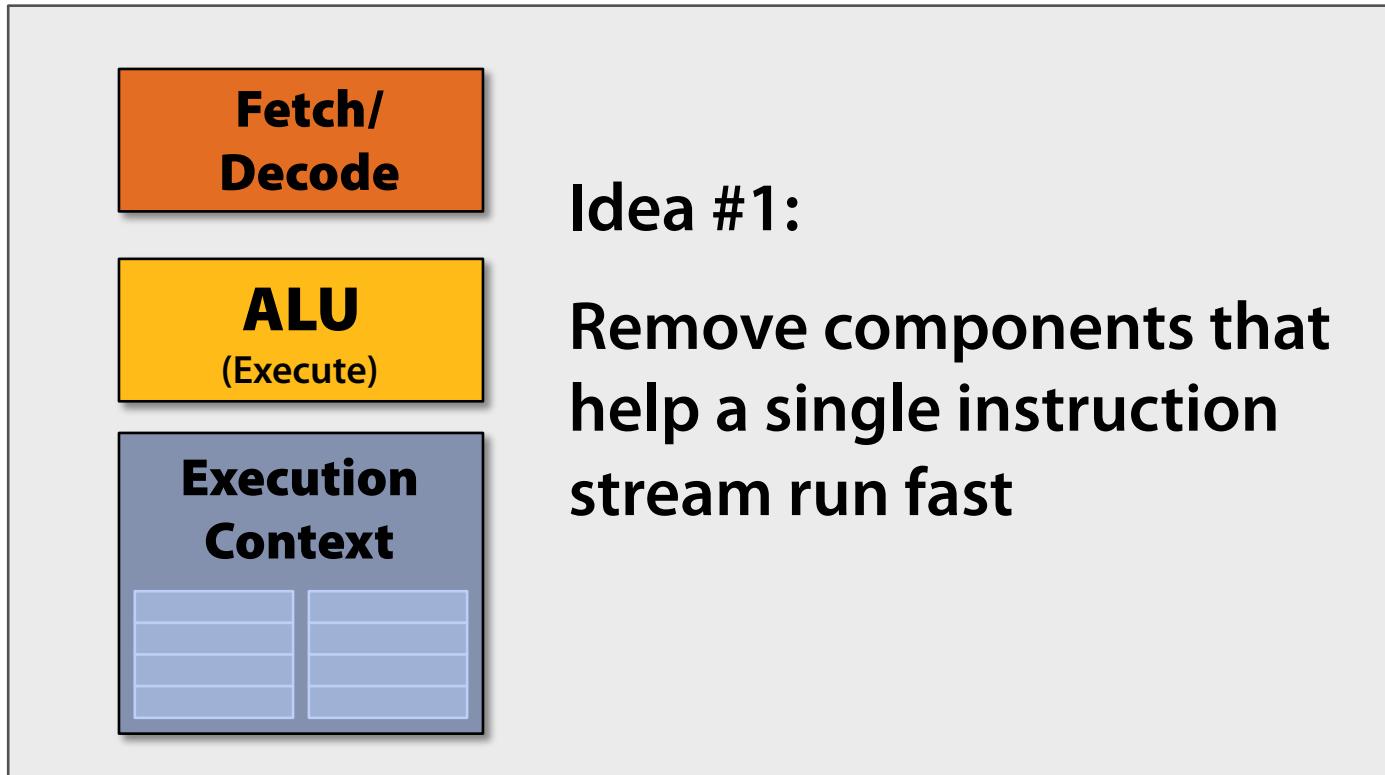
```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul r3, v0, cb0[0]  
madd r3, v1, cb0[1], r3  
madd r3, v2, cb0[2], r3  
clmp r3, r3, 1(0.0), 1(1.0)  
mul o0, r0, r3  
mul o1, r1, r3  
mul o2, r2, r3  
mov o3, 1(1.0)
```



CPU-“style” cores



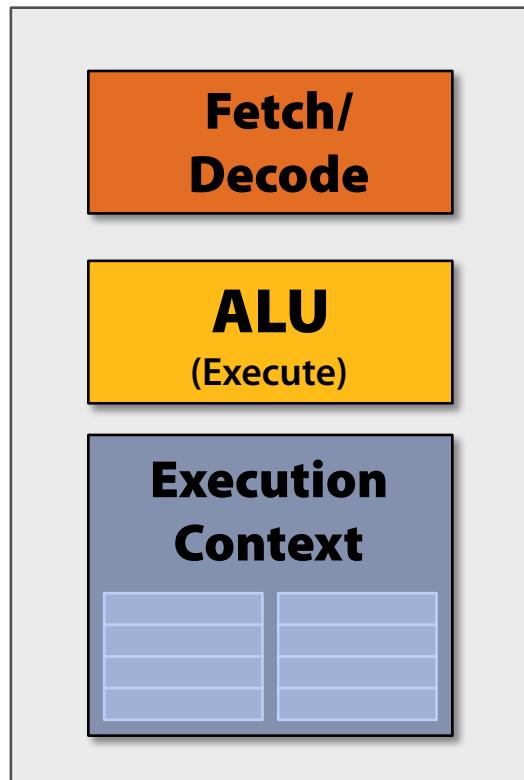
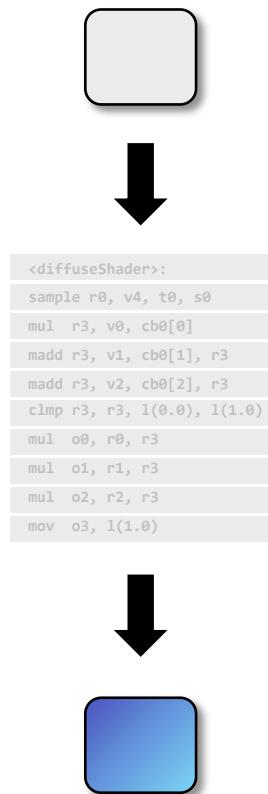
Slimming down



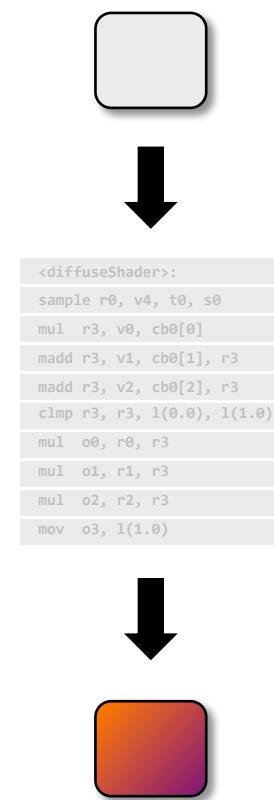
Two cores (two fragments in parallel)



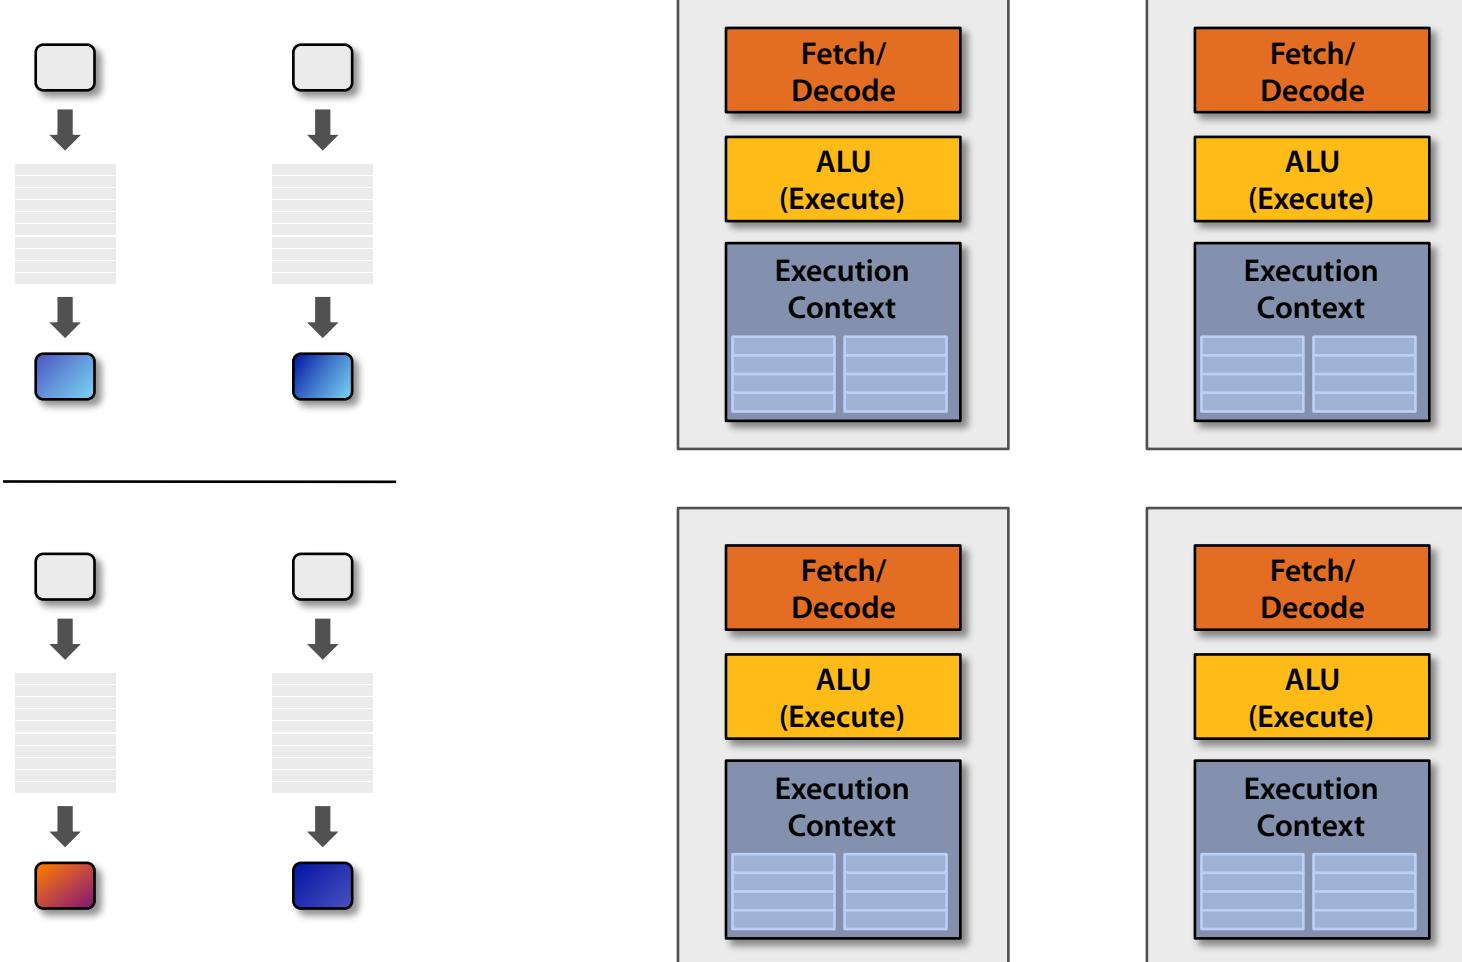
fragment 1



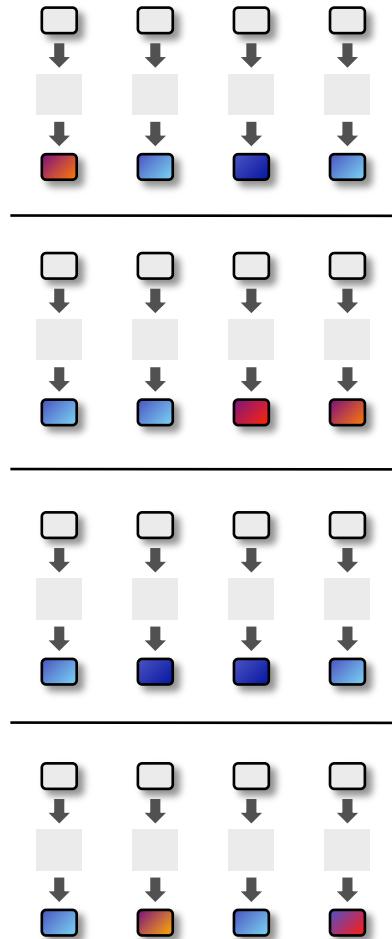
fragment 2



Four cores (four fragments in parallel)

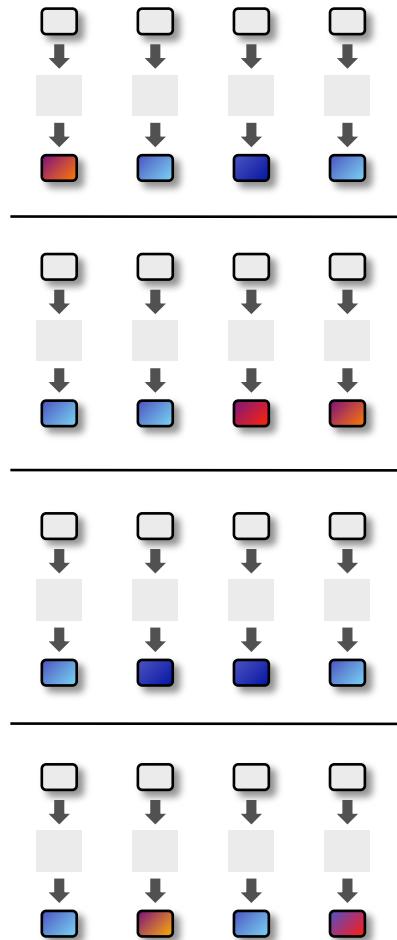


Sixteen cores (sixteen fragments in parallel)



16 cores = 16 simultaneous instruction streams

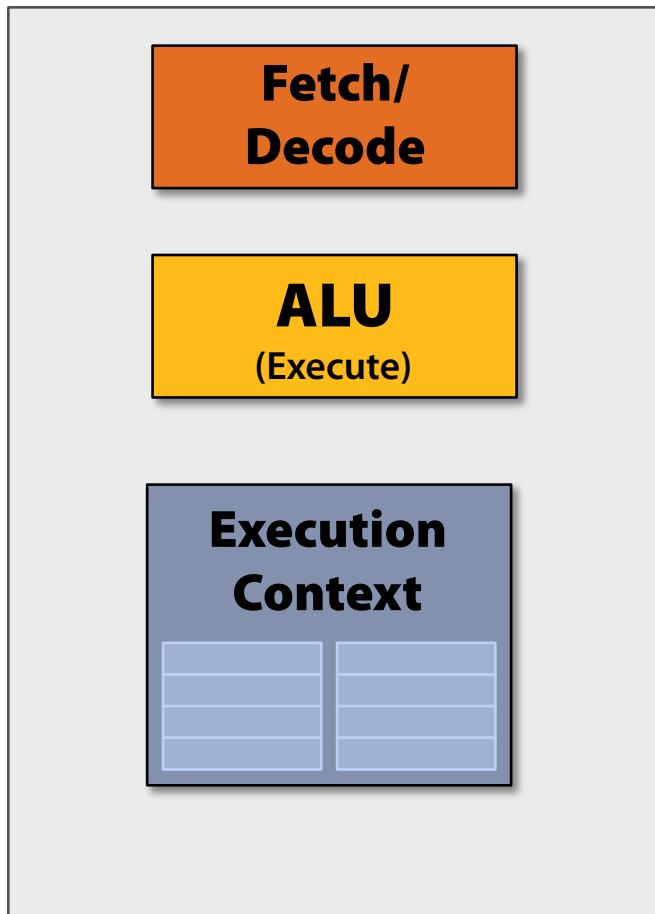
Instruction stream coherence



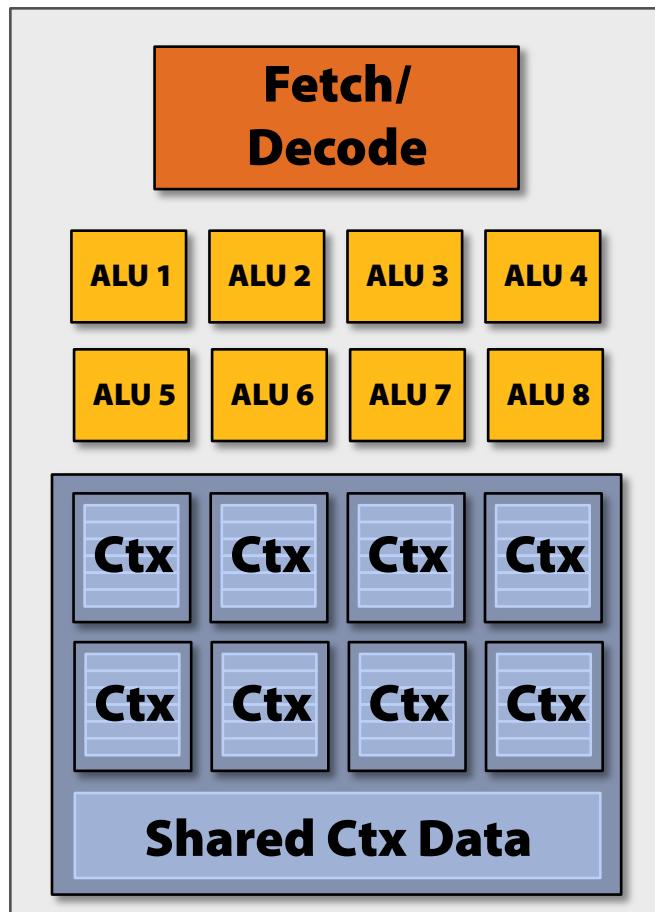
But... many fragments *should* be able to share an instruction stream!

```
<diffuseShader>:
sample r0, v4, t0, s0
mul r3, v0, cb0[0]
madd r3, v1, cb0[1], r3
madd r3, v2, cb0[2], r3
clmp r3, r3, 1(0.0), 1(1.0)
mul o0, r0, r3
mul o1, r1, r3
mul o2, r2, r3
mov o3, 1(1.0)
```

Recall: simple processing core



Add ALUs

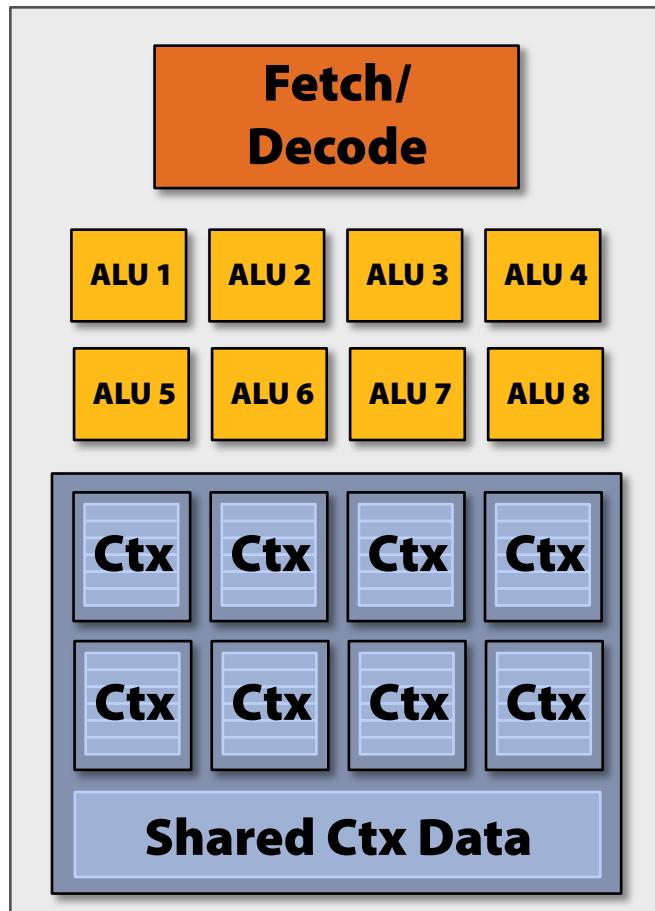


Idea #2:

Amortize cost/complexity of managing an instruction stream across many ALUs

SIMD processing

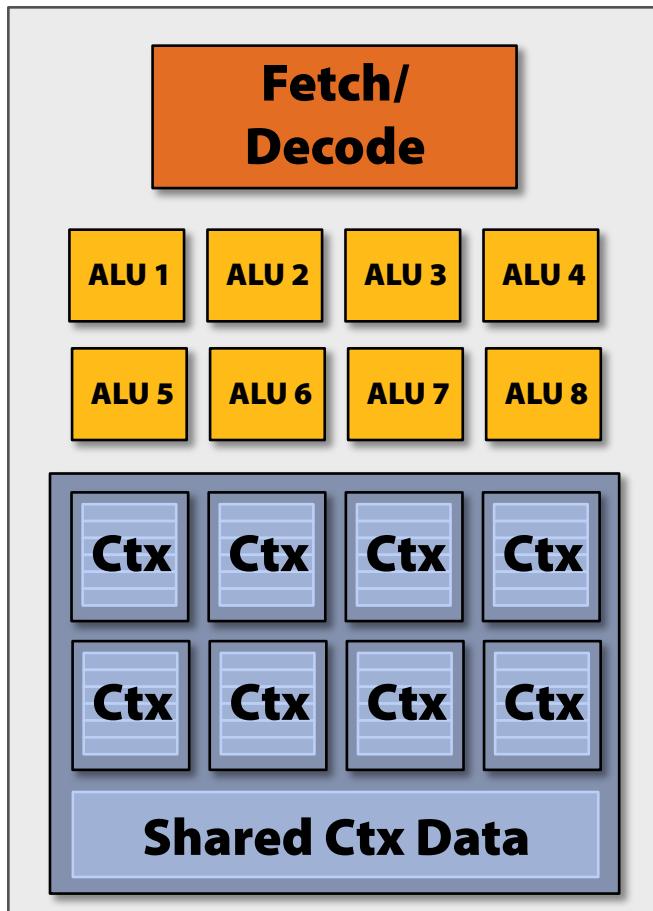
Modifying the shader



```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul r3, v0, cb0[0]  
madd r3, v1, cb0[1], r3  
madd r3, v2, cb0[2], r3  
clmp r3, r3, 1(0.0), 1(1.0)  
mul o0, r0, r3  
mul o1, r1, r3  
mul o2, r2, r3  
mov o3, 1(1.0)
```

Original compiled shader:
Processes one fragment
using scalar ops on scalar
registers

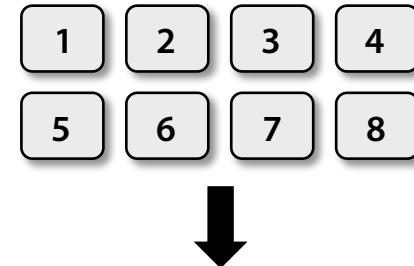
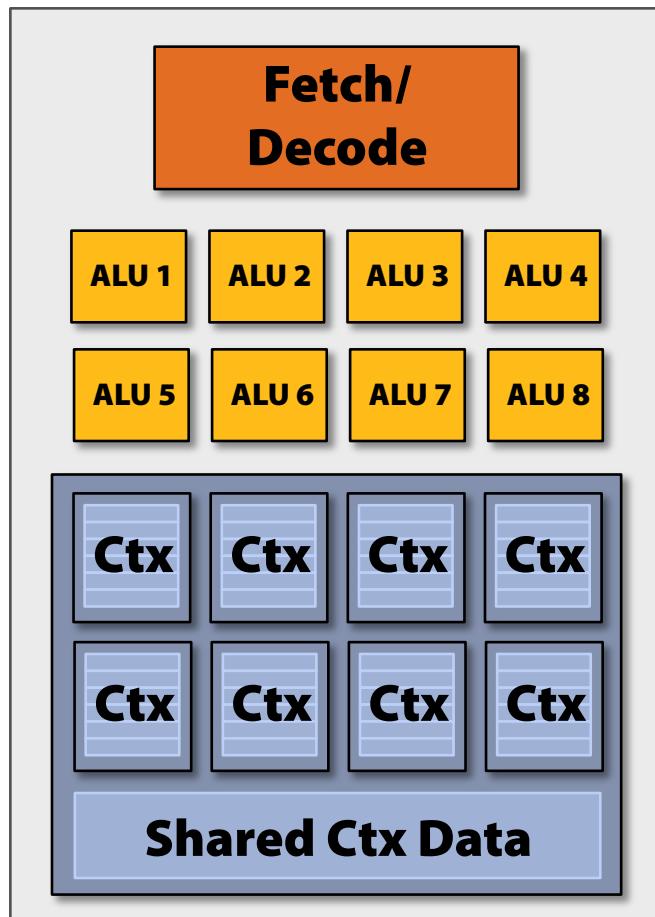
Modifying the shader



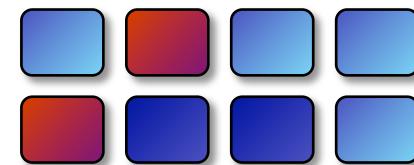
```
<VEC8_diffuseShader>:  
VEC8_sample vec_r0, vec_v4, t0, vec_s0  
VEC8_mul  vec_r3, vec_v0, cb0[0]  
VEC8_madd vec_r3, vec_v1, cb0[1], vec_r3  
VEC8_madd vec_r3, vec_v2, cb0[2], vec_r3  
VEC8_clmp vec_r3, vec_r3, 1(0.0), 1(1.0)  
VEC8_mul  vec_o0, vec_r0, vec_r3  
VEC8_mul  vec_o1, vec_r1, vec_r3  
VEC8_mul  vec_o2, vec_r2, vec_r3  
VEC8_mov  vec_o3, 1(1.0)
```

New compiled shader:
Processes 8 fragments
using vector ops on vector
registers

Modifying the shader



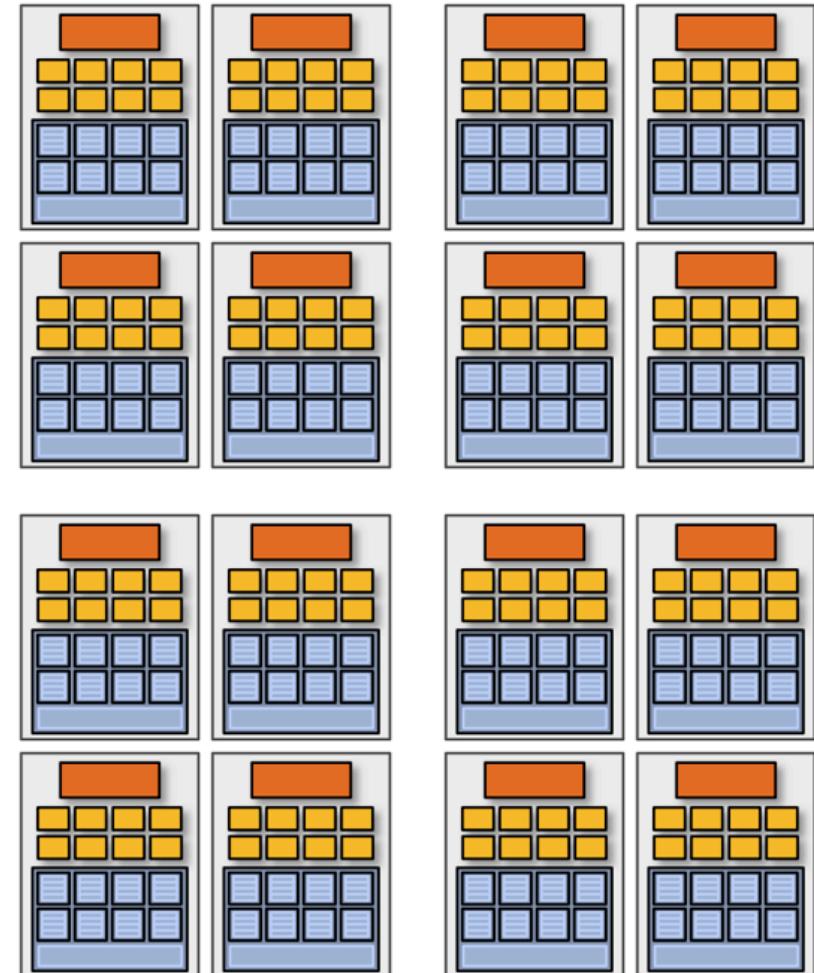
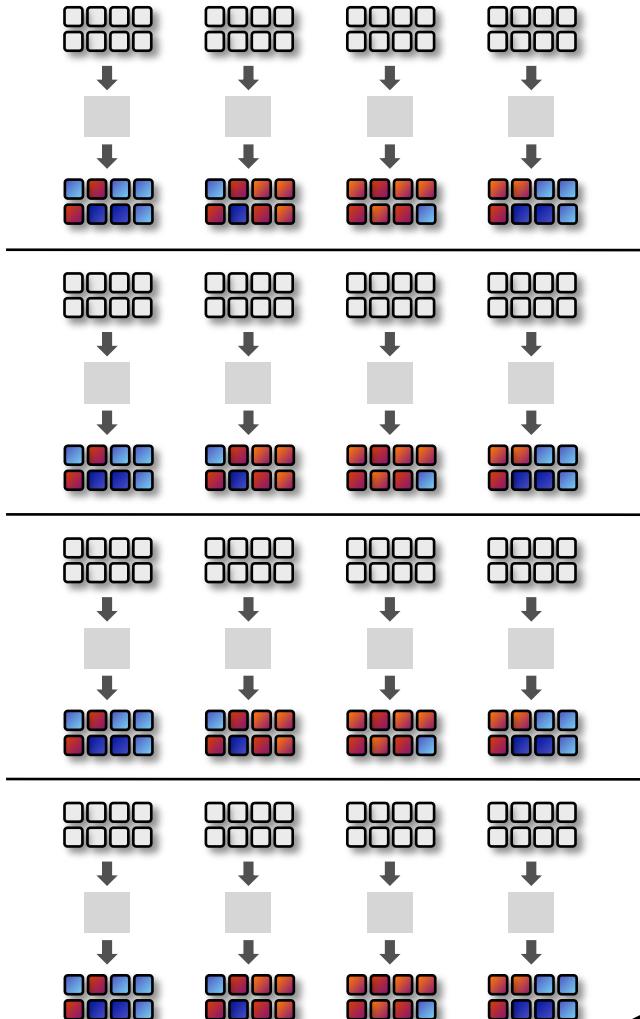
```
<VEC8_diffuseShader>:
VEC8_sample vec_r0, vec_v4, t0, vec_s0
VEC8_mul  vec_r3, vec_v0, cb0[0]
VEC8_madd vec_r3, vec_v1, cb0[1], vec_r3
VEC8_madd vec_r3, vec_v2, cb0[2], vec_r3
VEC8_clmp vec_r3, vec_r3, 1(0.0), 1(1.0)
VEC8_mul  vec_o0, vec_r0, vec_r3
VEC8_mul  vec_o1, vec_r1, vec_r3
VEC8_mul  vec_o2, vec_r2, vec_r3
VEC8_mov  vec_o3, 1(1.0)
```





SIGGRAPH2008

128 fragments in parallel



**16 cores = 128 ALUs
= 16 simultaneous instruction streams**



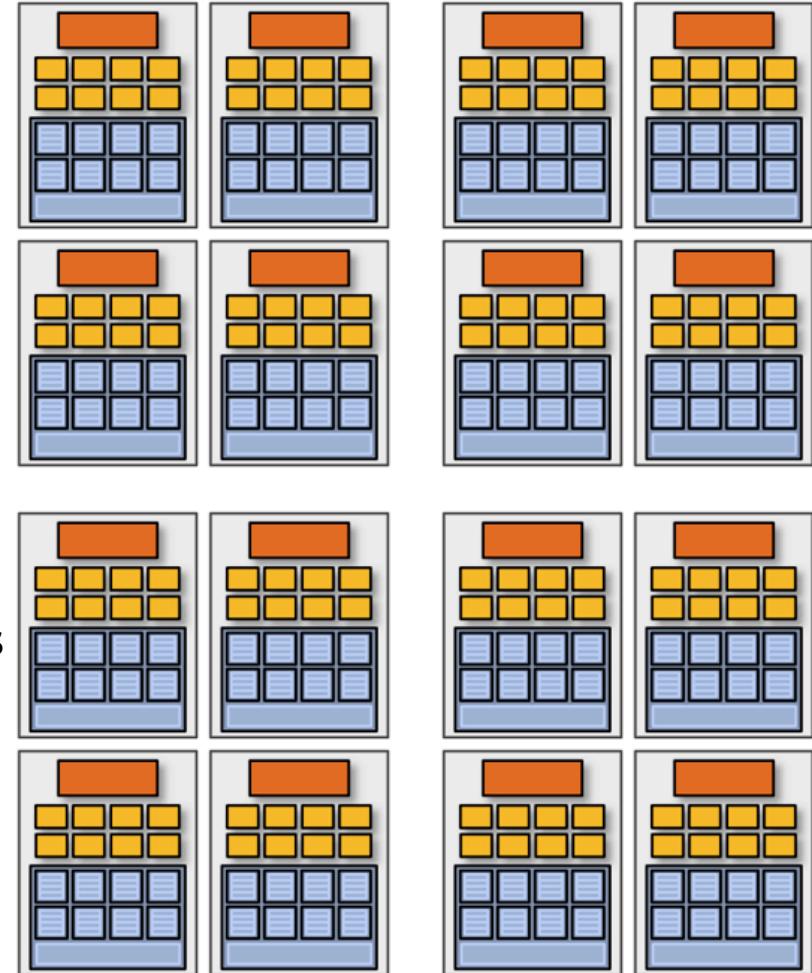
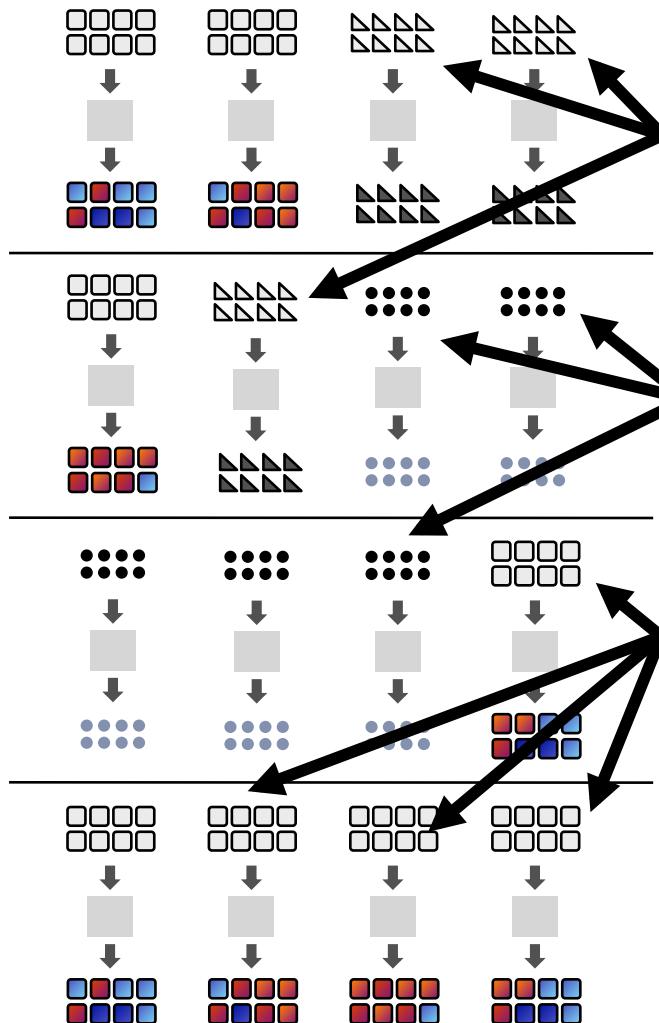
SIGGRAPH2008

128 [

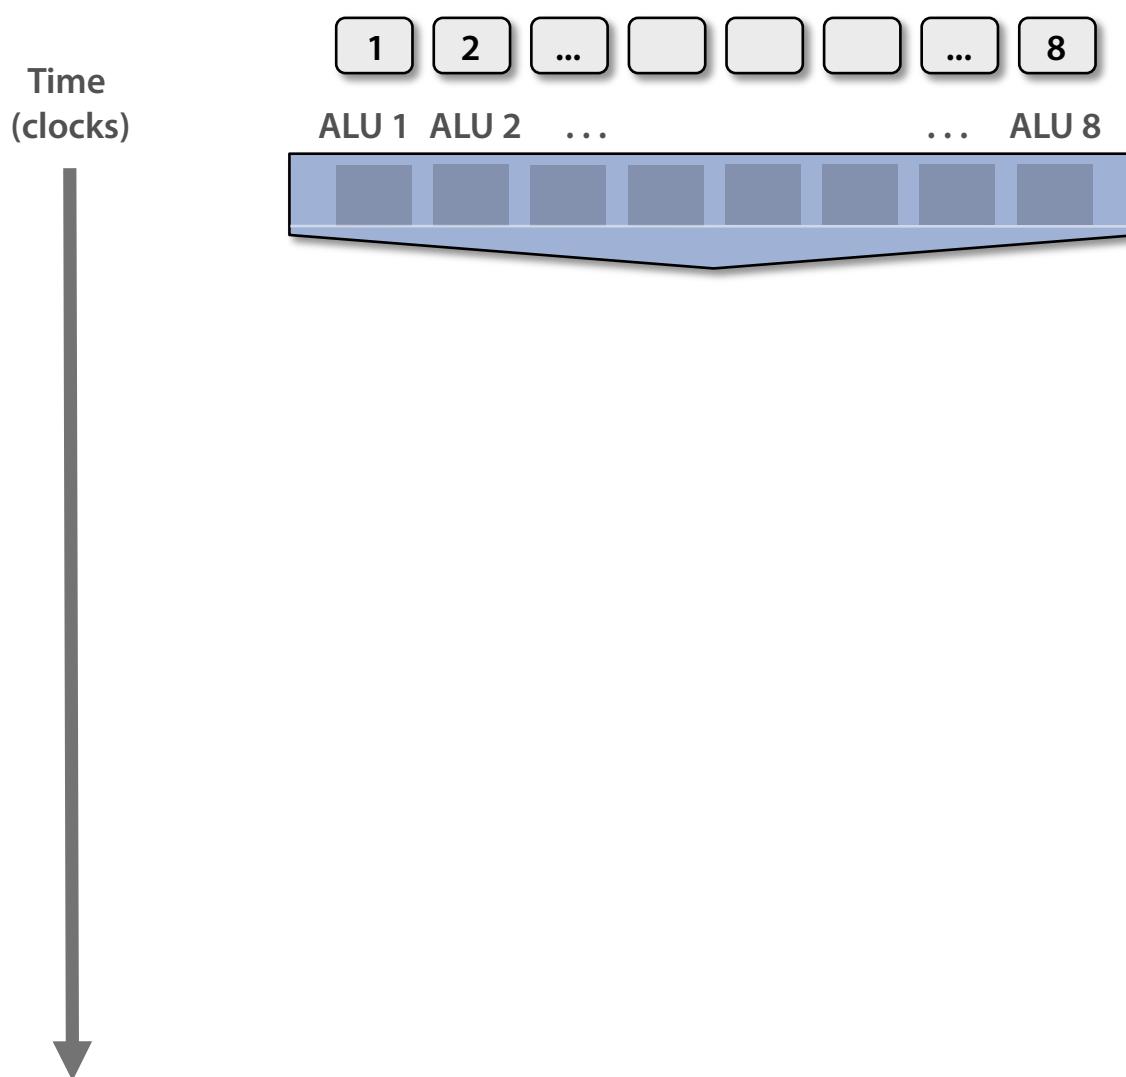
vertices / fragments
primitives
CUDA threads
OpenCL work items

] in parallel

compute shader threads



But what about branches?

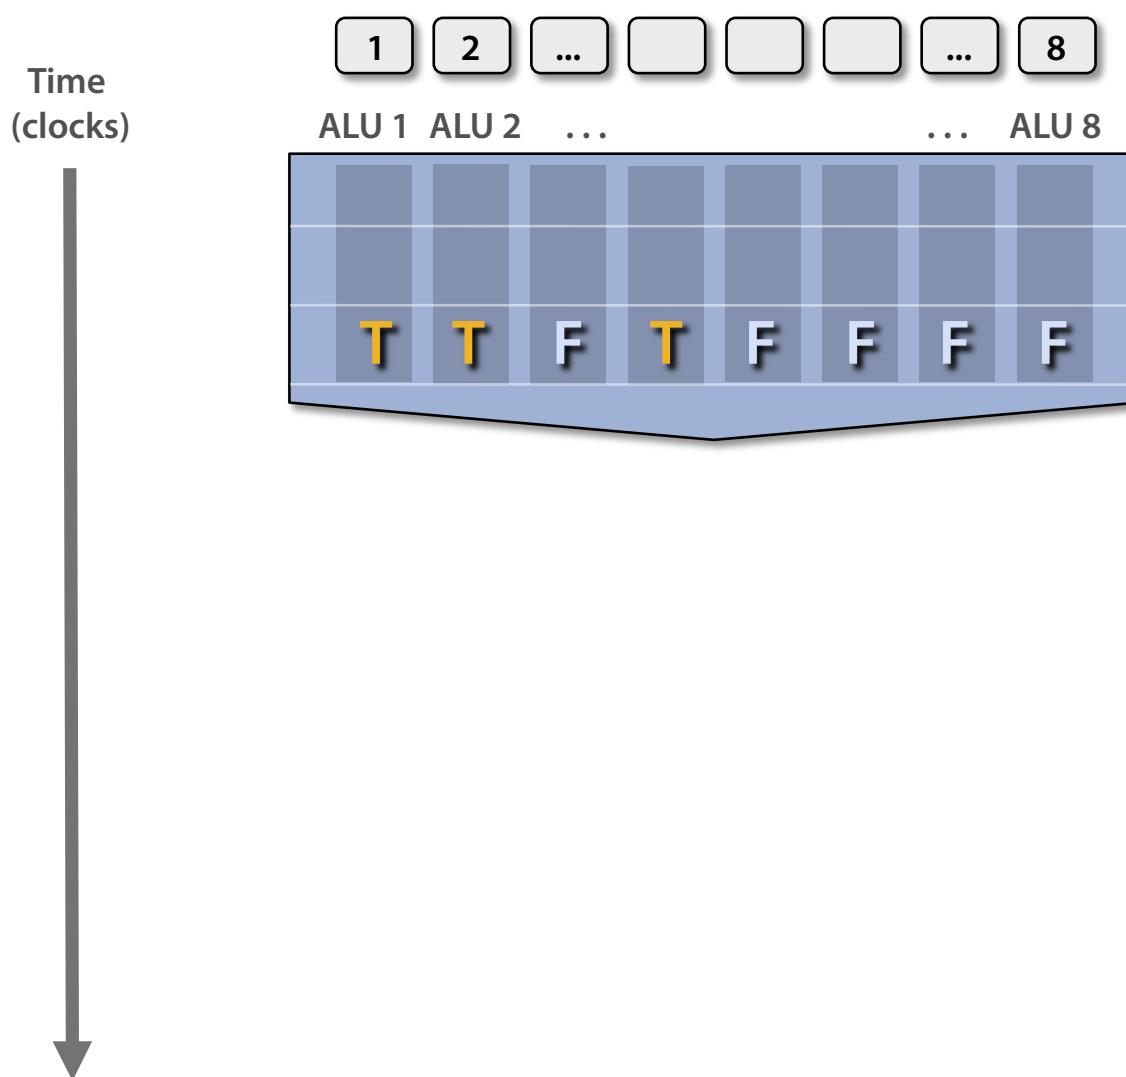


<unconditional
shader code>

```
if (x > 0) {  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y + Ka;  
} else {  
    x = 0;  
    refl = Ka;  
}
```

<resume unconditional
shader code>

But what about branches?

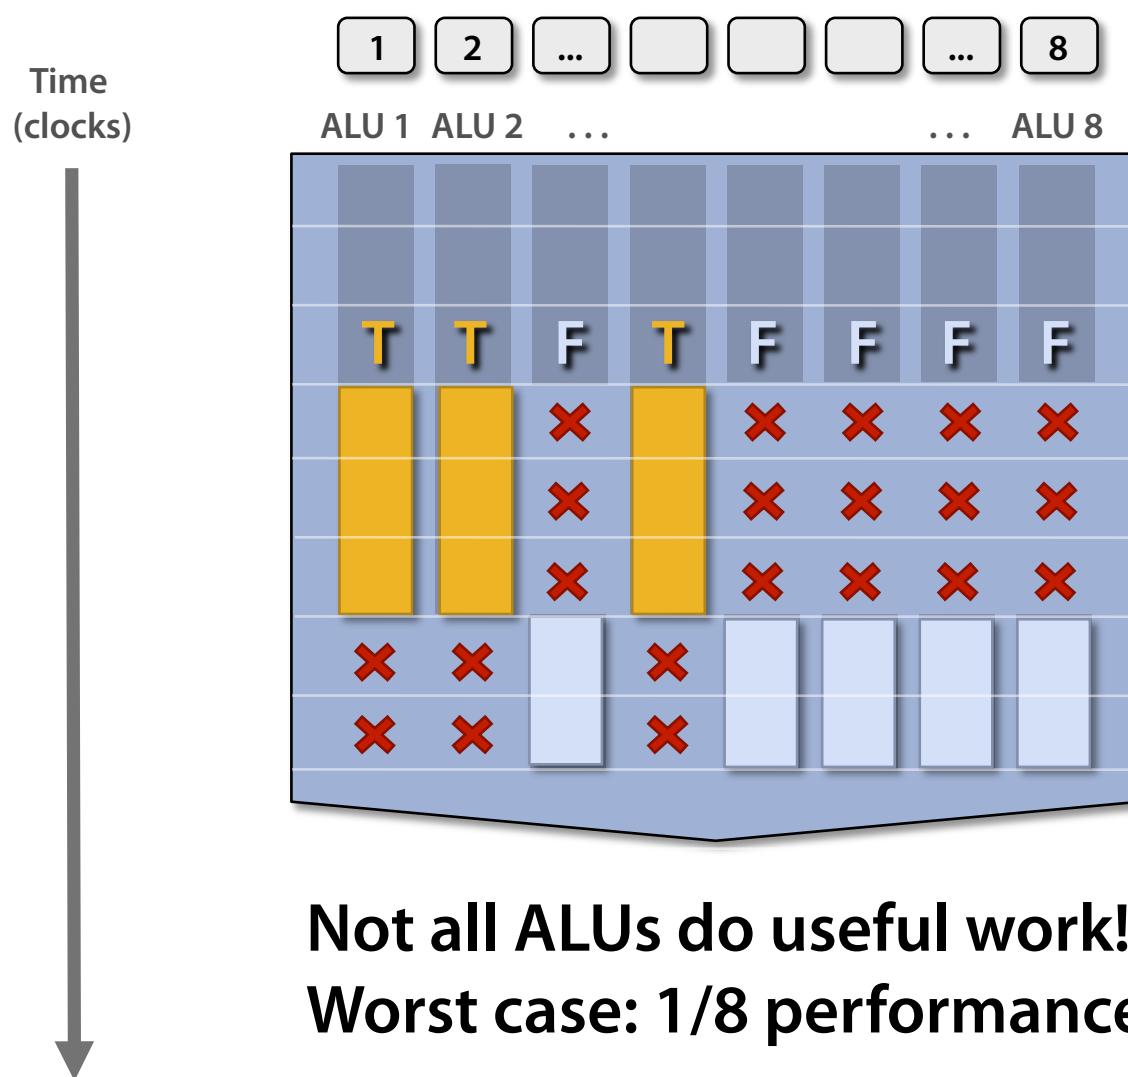


<unconditional
shader code>

```
if (x > 0) {  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y + Ka;  
} else {  
    x = 0;  
    refl = Ka;  
}
```

<resume unconditional
shader code>

But what about branches?

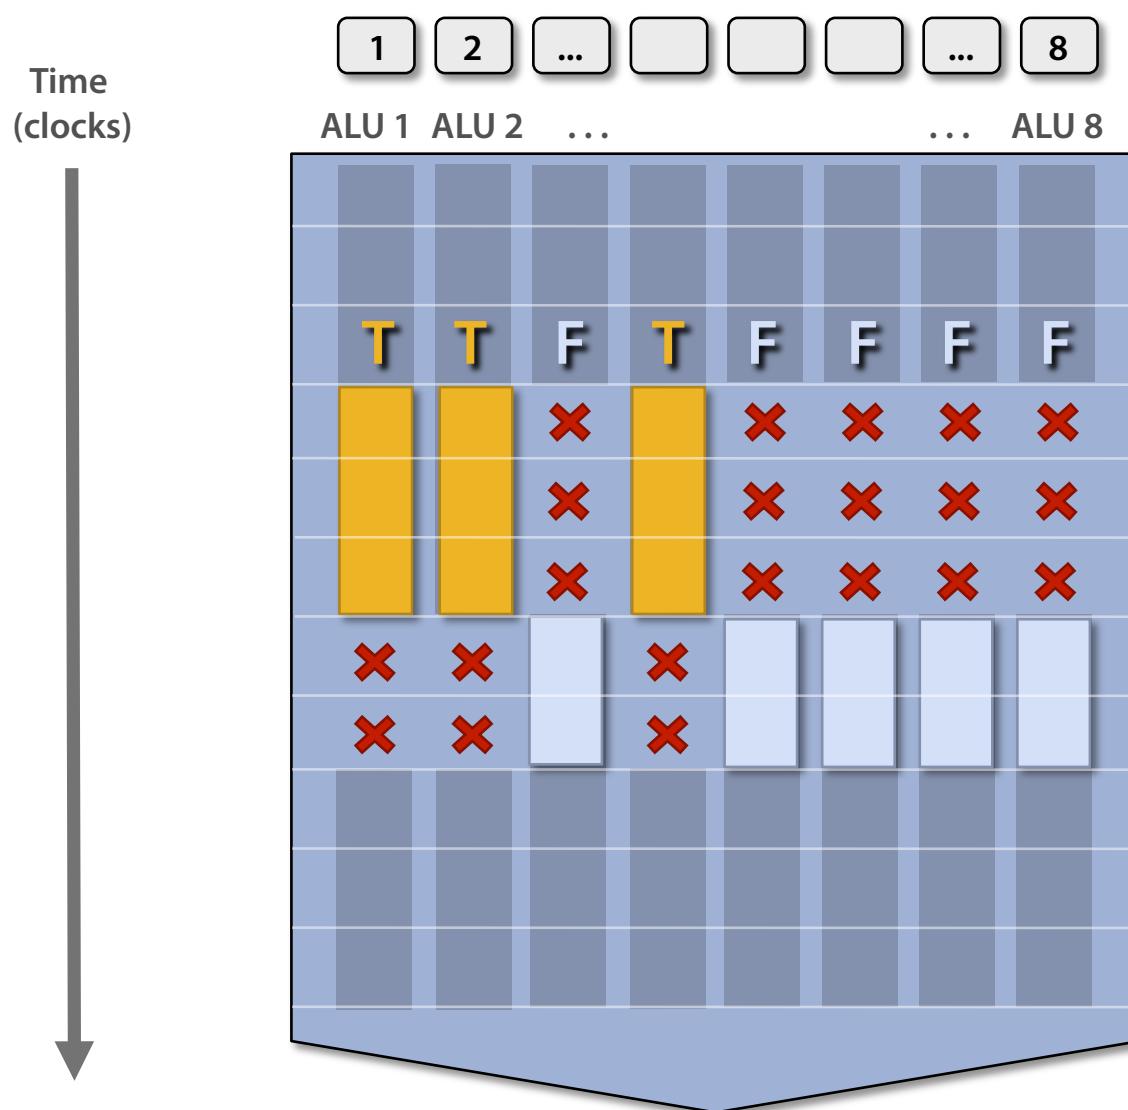


<unconditional shader code>

```
if (x > 0) {
    y = pow(x, exp);
    y *= Ks;
    refl = y + Ka;
} else {
    x = 0;
    refl = Ka;
}
```

<resume unconditional shader code>

But what about branches?



<unconditional shader code>

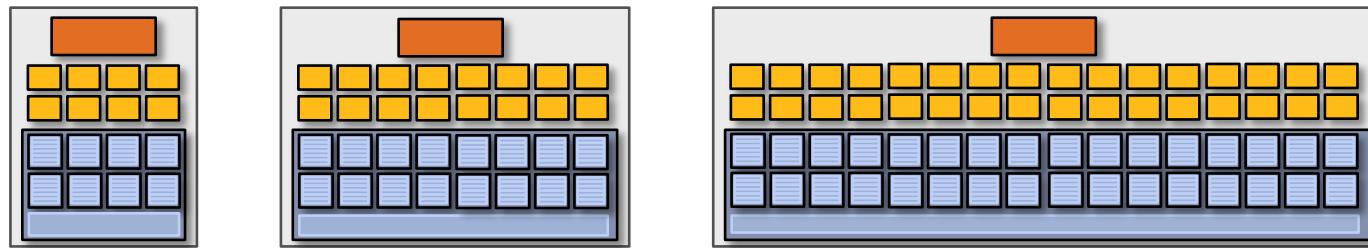
```
if (x > 0) {
    y = pow(x, exp);
    y *= Ks;
    refl = y + Ka;
} else {
    x = 0;
    refl = Ka;
}
```

<resume unconditional shader code>

Clarification

SIMD processing does not imply SIMD instructions

- Option 1: Explicit vector instructions
 - Intel/AMD x86 SSE, Intel Larrabee
- Option 2: Scalar instructions, implicit HW vectorization
 - HW determines instruction stream sharing across ALUs (amount of sharing hidden from software)
 - NVIDIA GeForce (“SIMT” warps), AMD Radeon architectures



In practice: 16 to 64 fragments share an instruction stream



SIGGRAPH2008

Stalls!

Stalls occur when a core cannot run the next instruction because of a dependency on a previous operation.

Texture access latency = 100's to 1000's of cycles

We've removed the fancy caches and logic that helps avoid stalls.



SIGGRAPH2008

But we have **LOTS** of independent fragments.

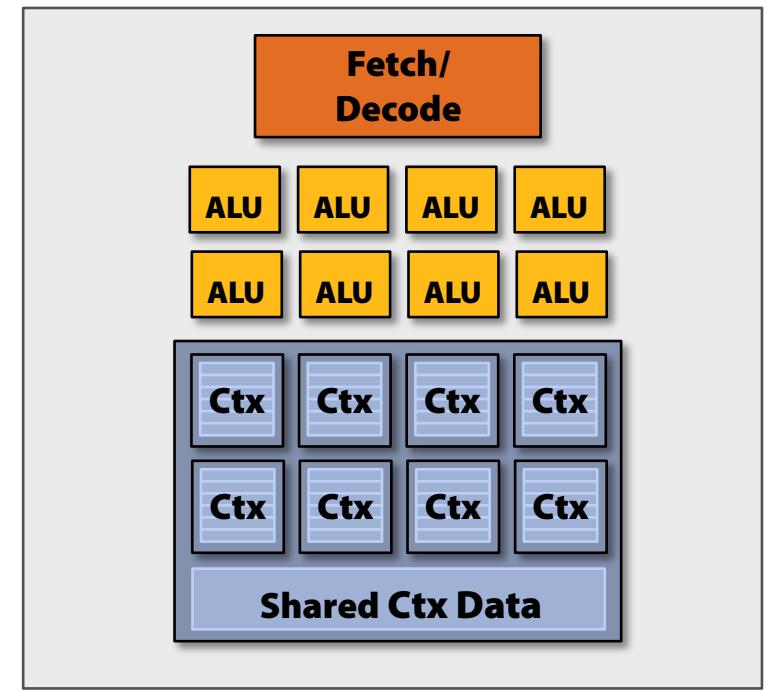
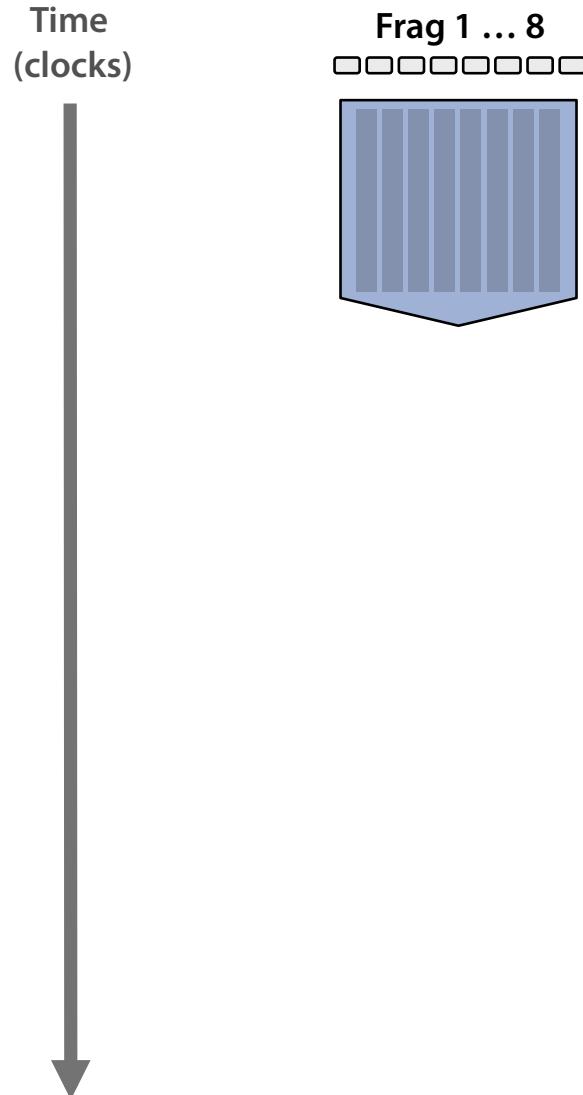
Idea #3:

Interleave processing of many fragments on a single core to avoid stalls caused by high latency operations.



SIGGRAPH2008

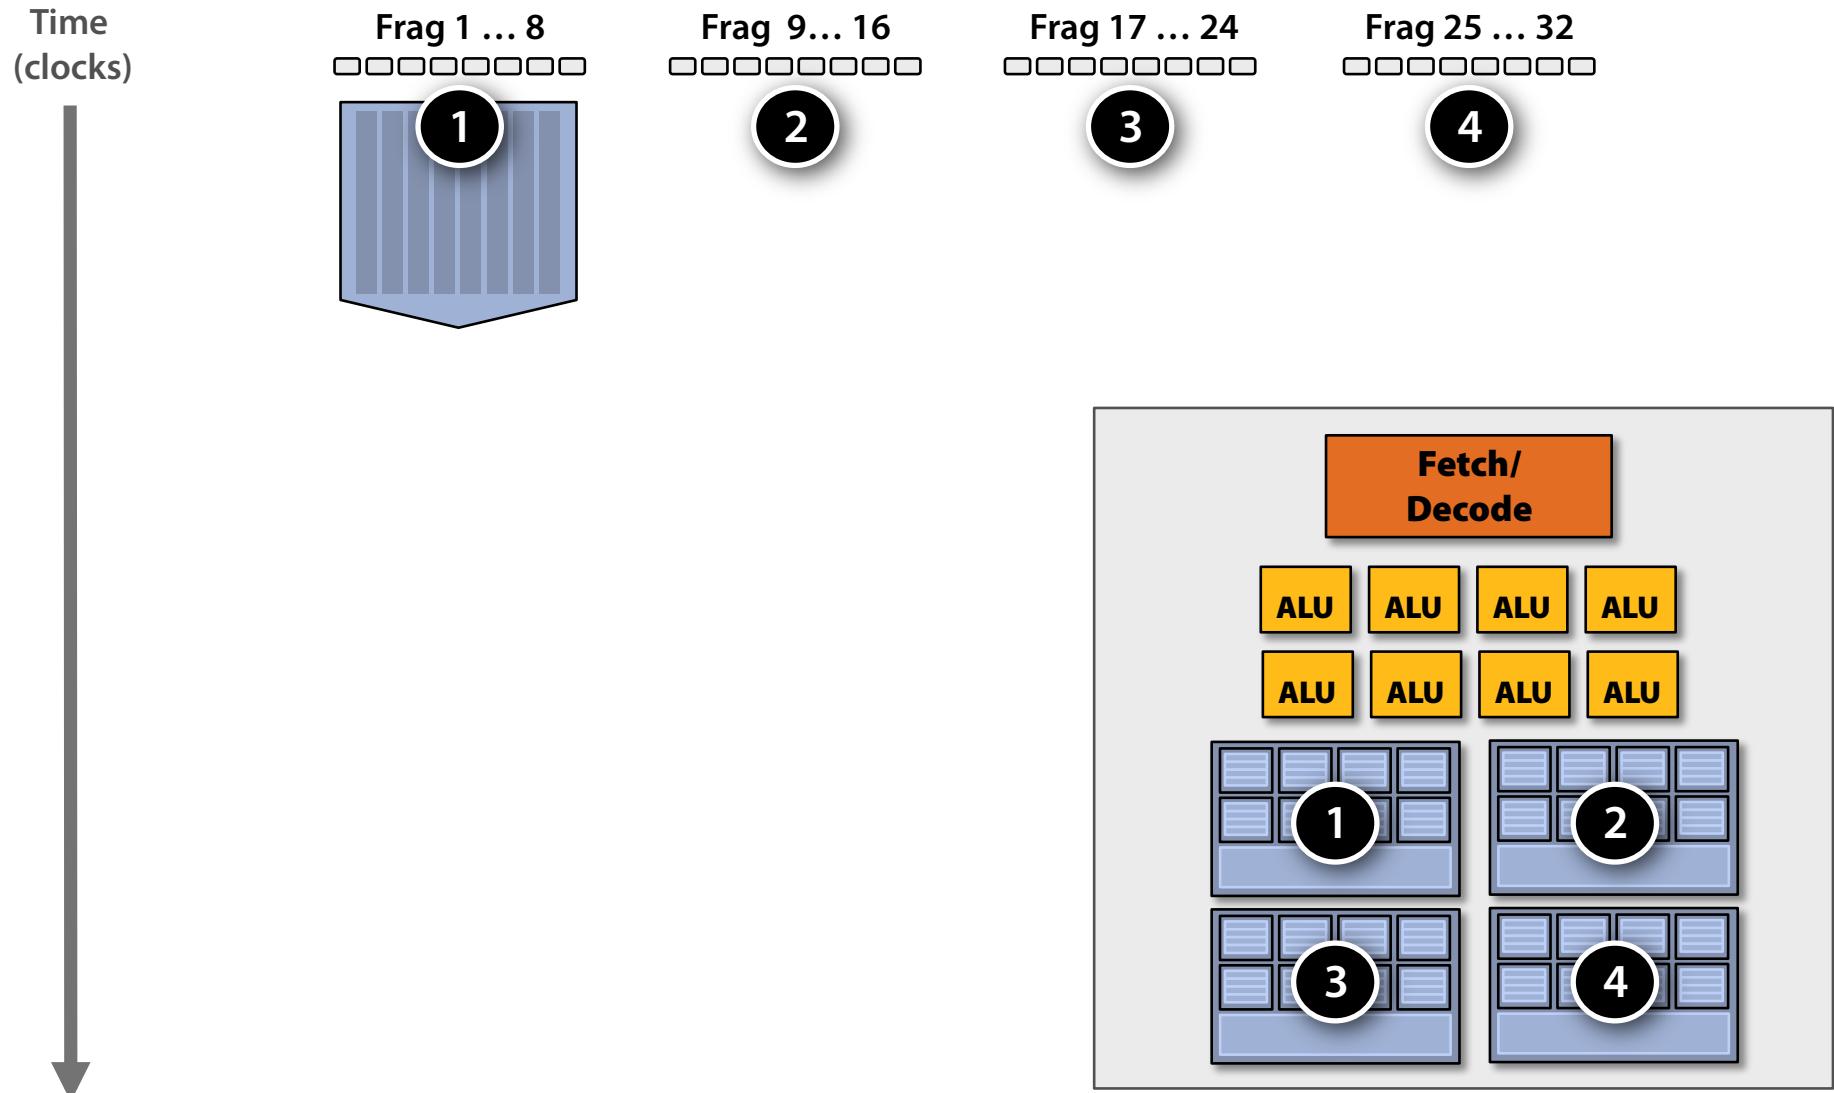
Hiding shader stalls





SIGGRAPH2008

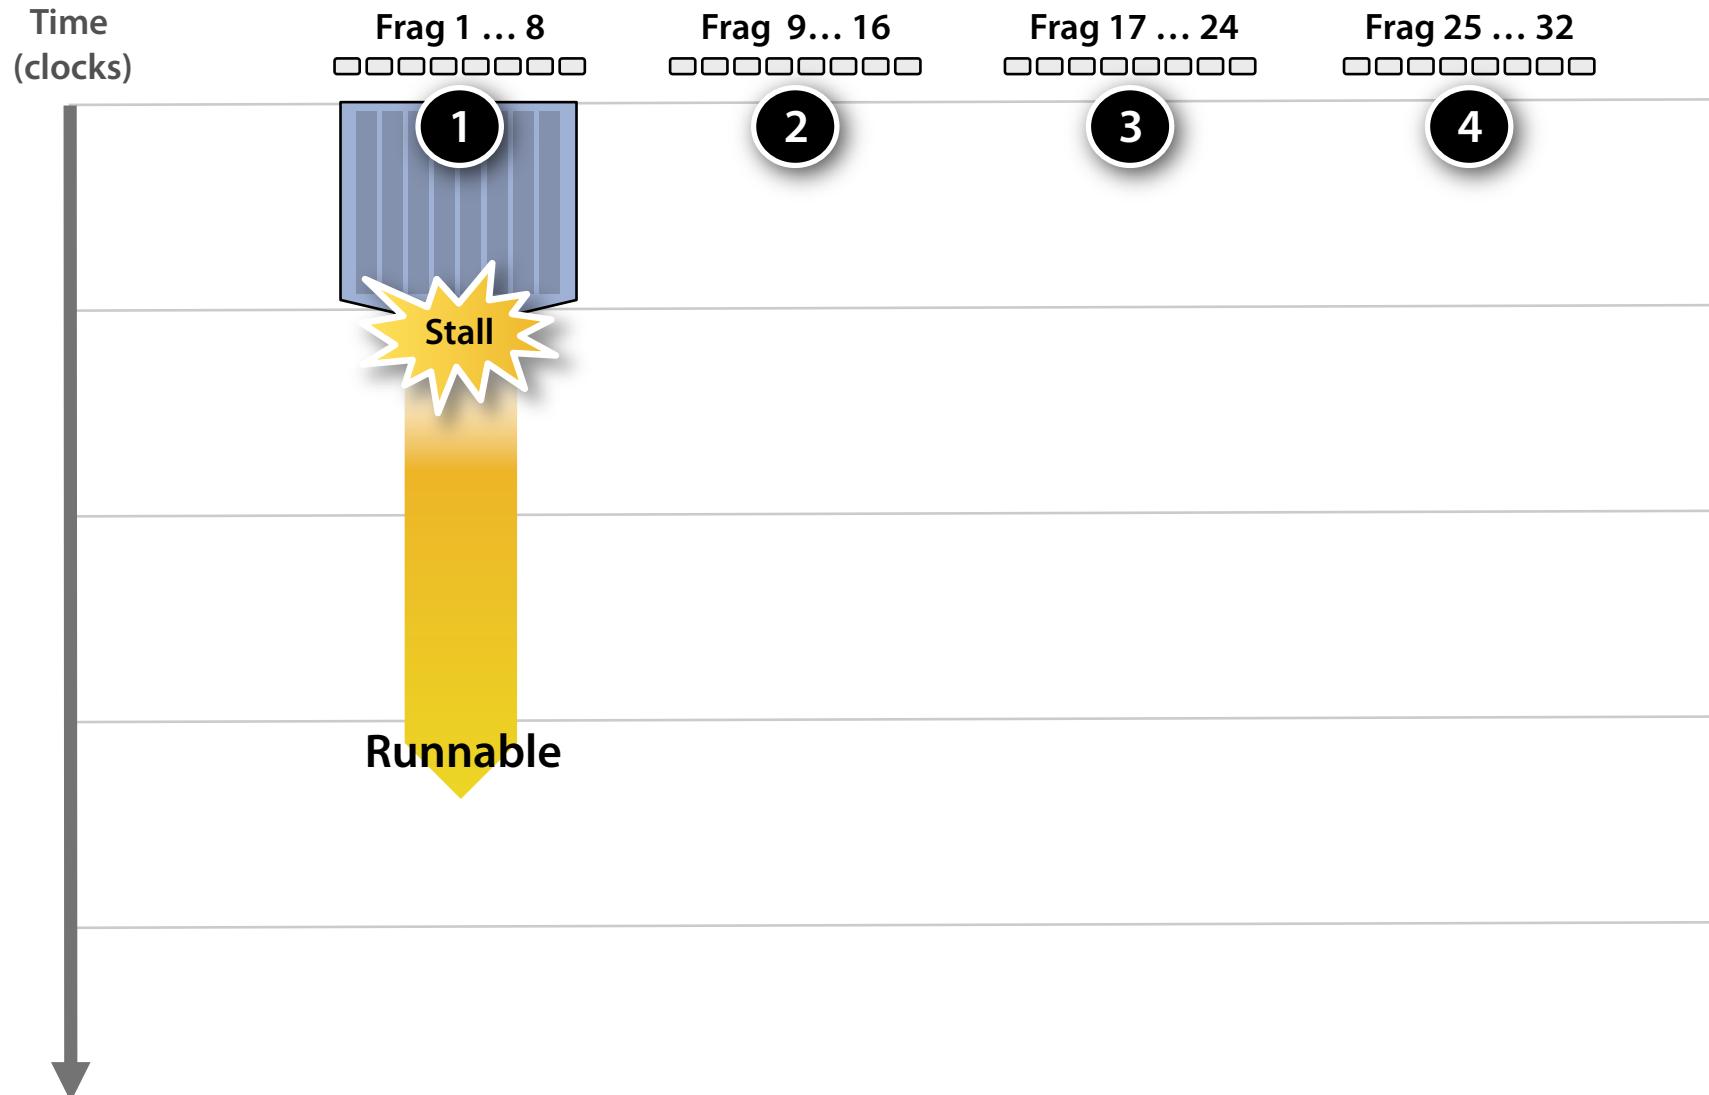
Hiding shader stalls





SIGGRAPH2008

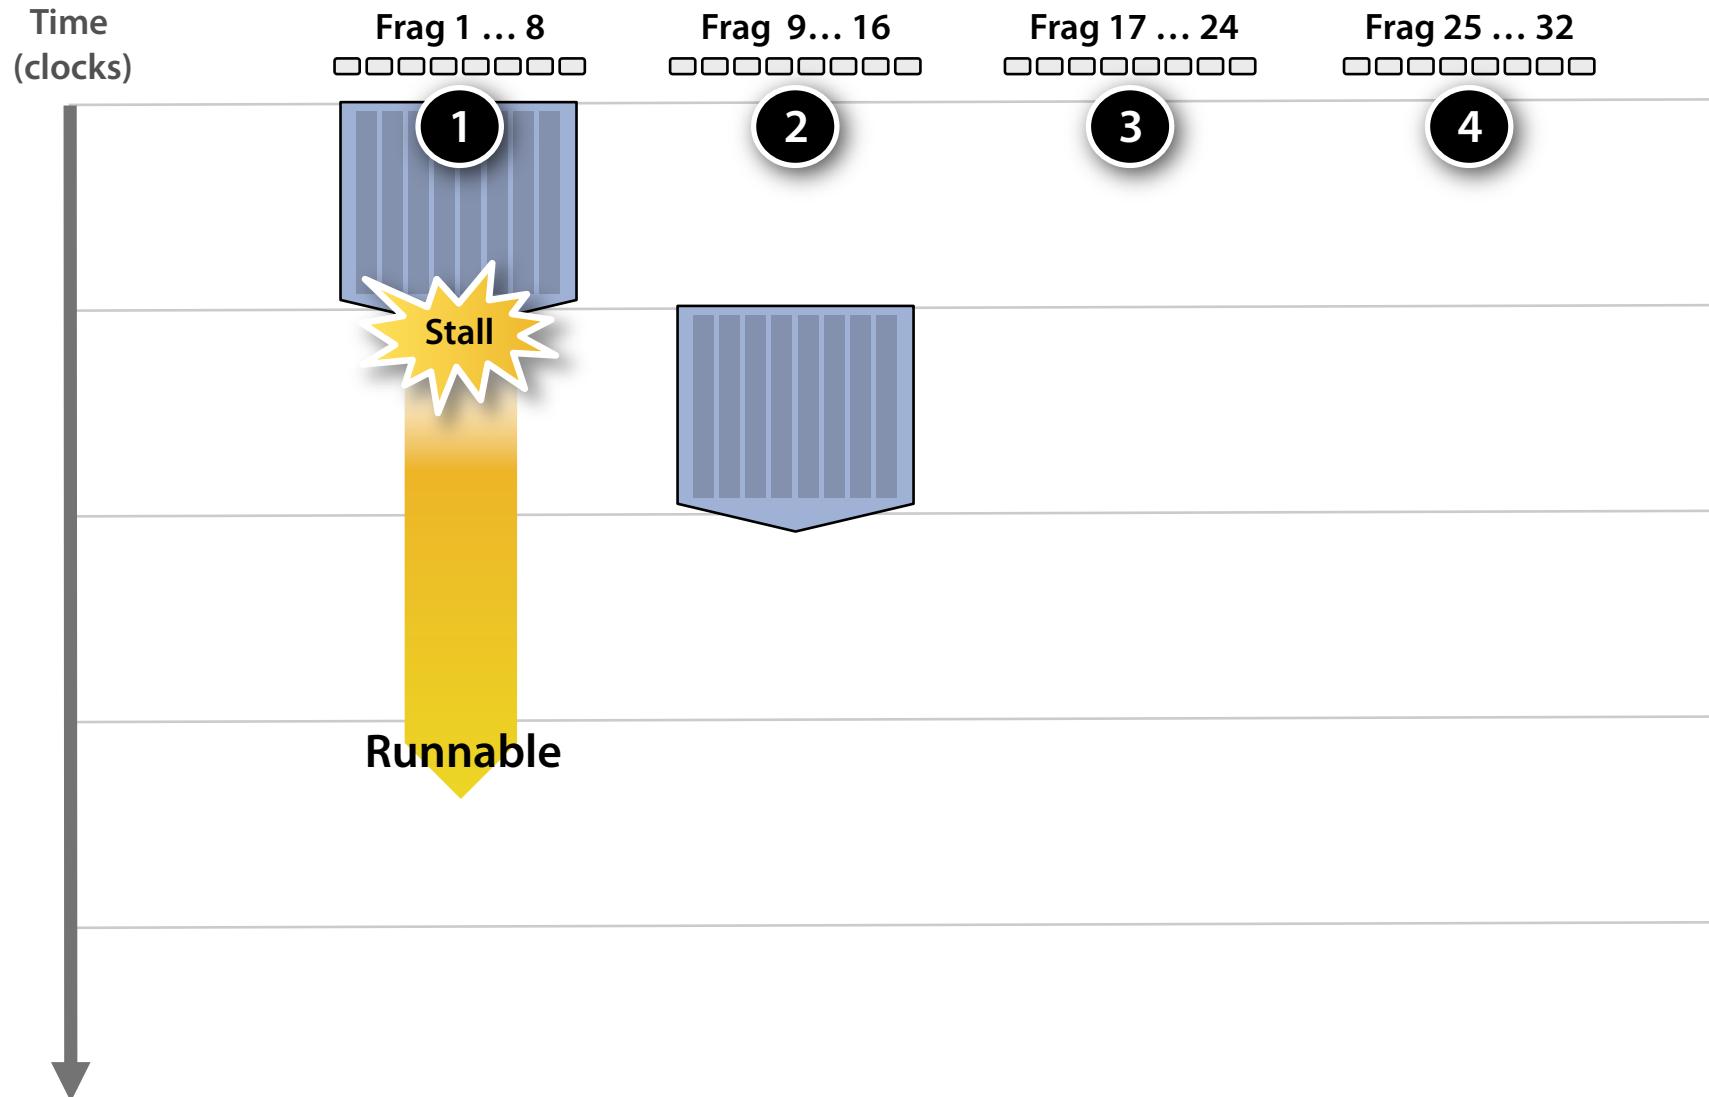
Hiding shader stalls



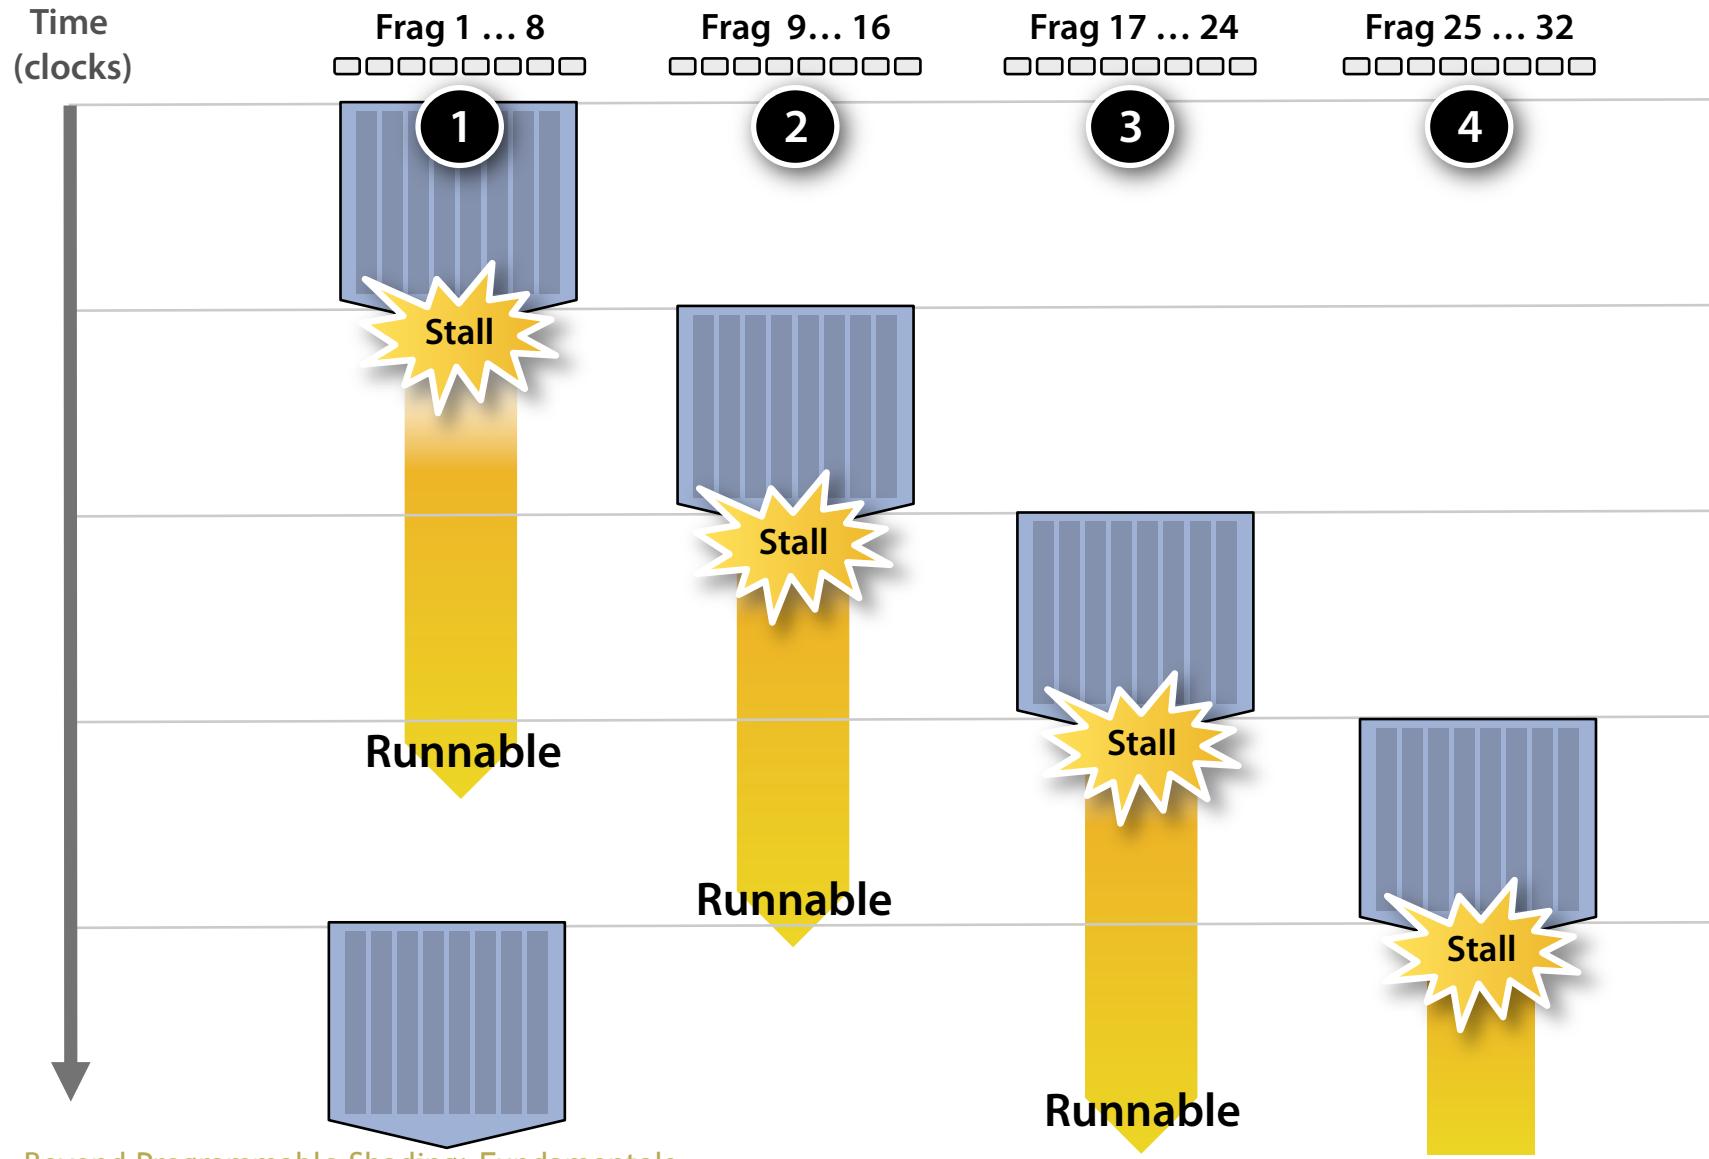


SIGGRAPH2008

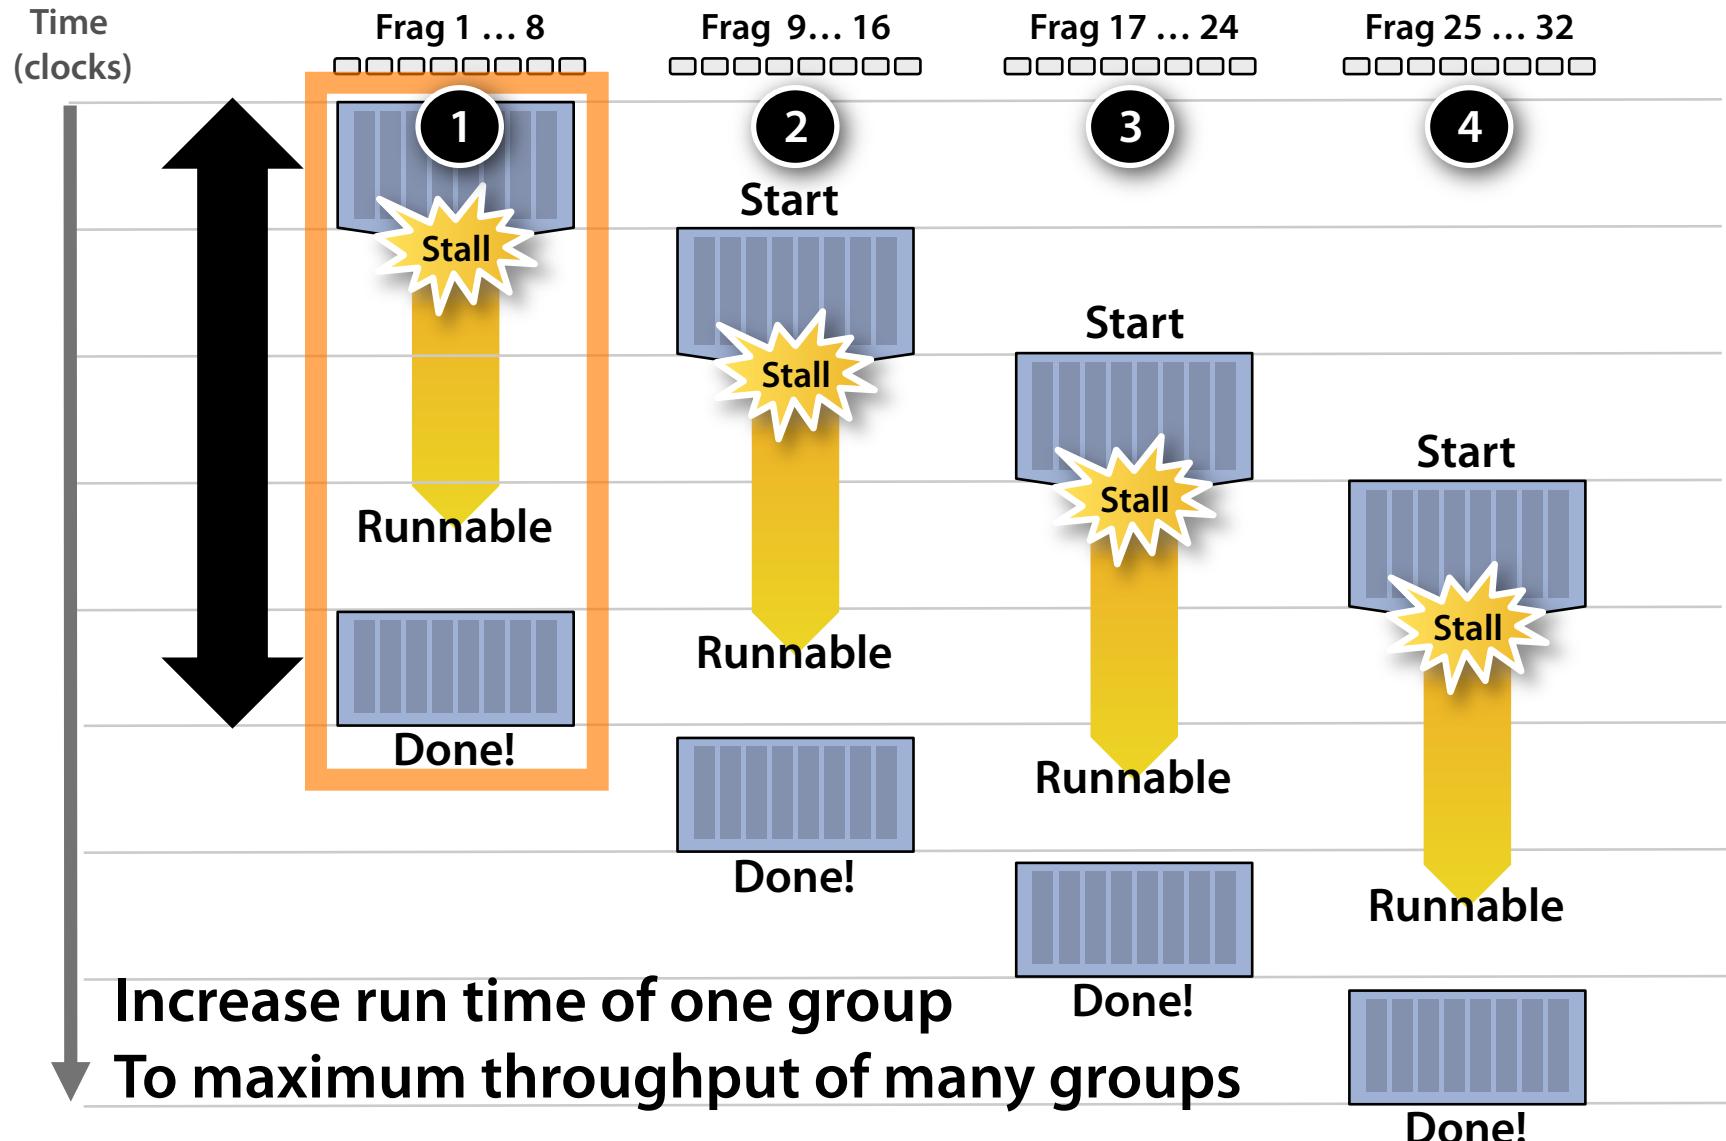
Hiding shader stalls



Hiding shader stalls



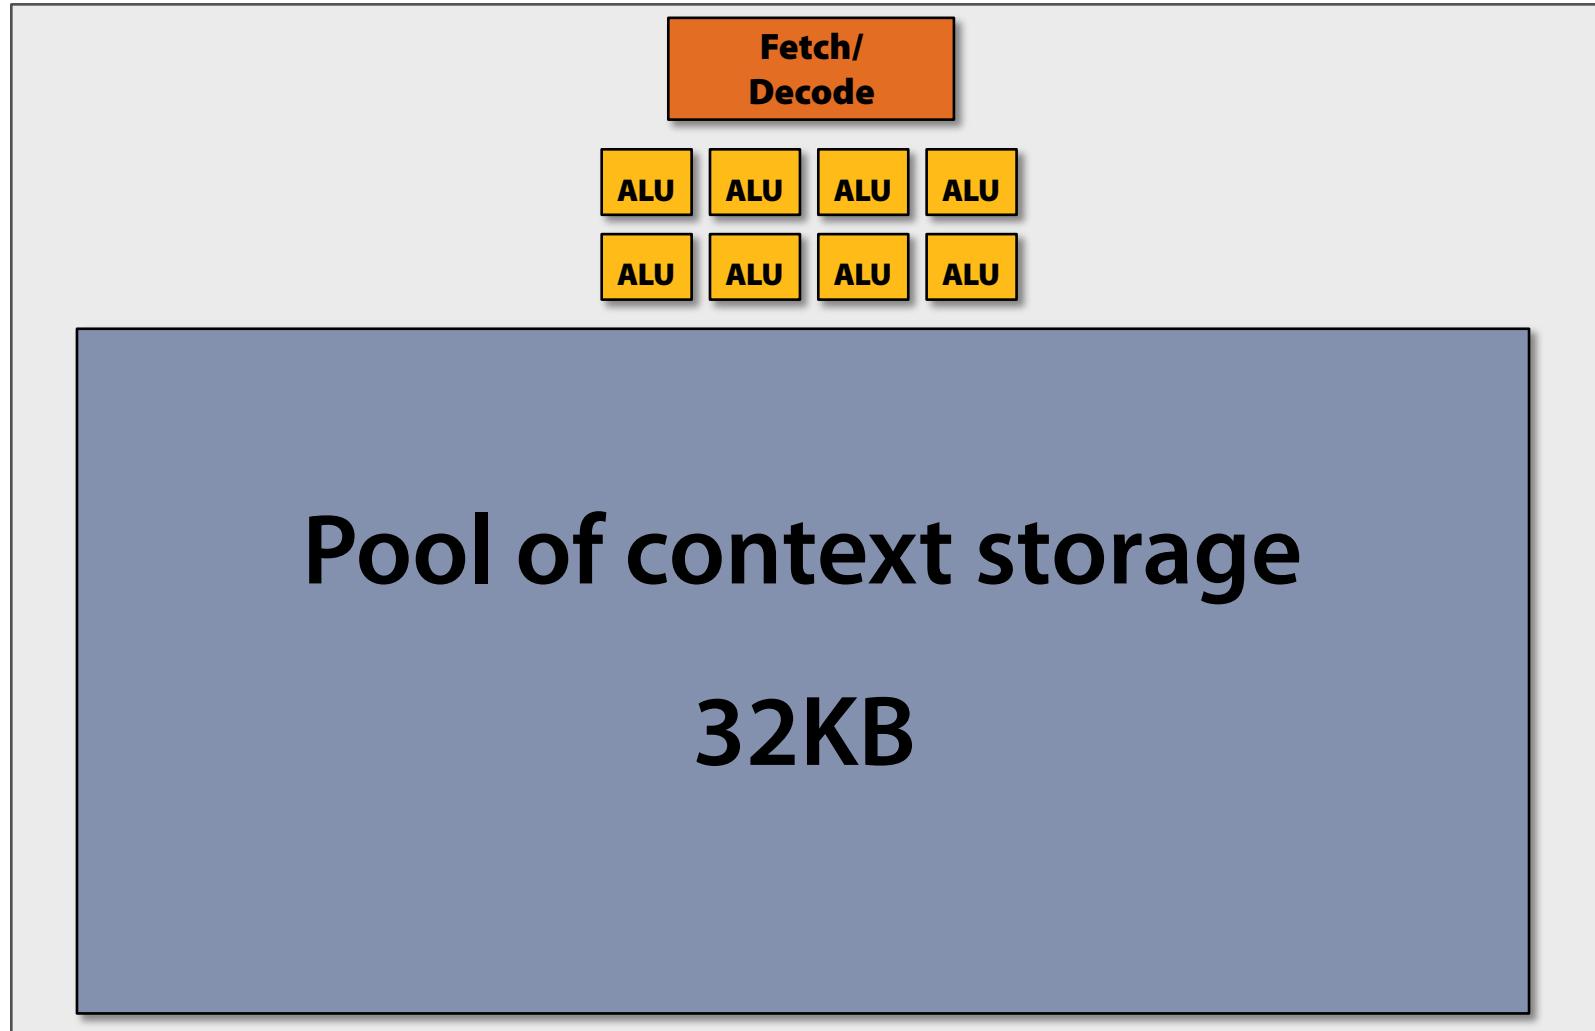
Throughput!





SIGGRAPH2008

Storing contexts

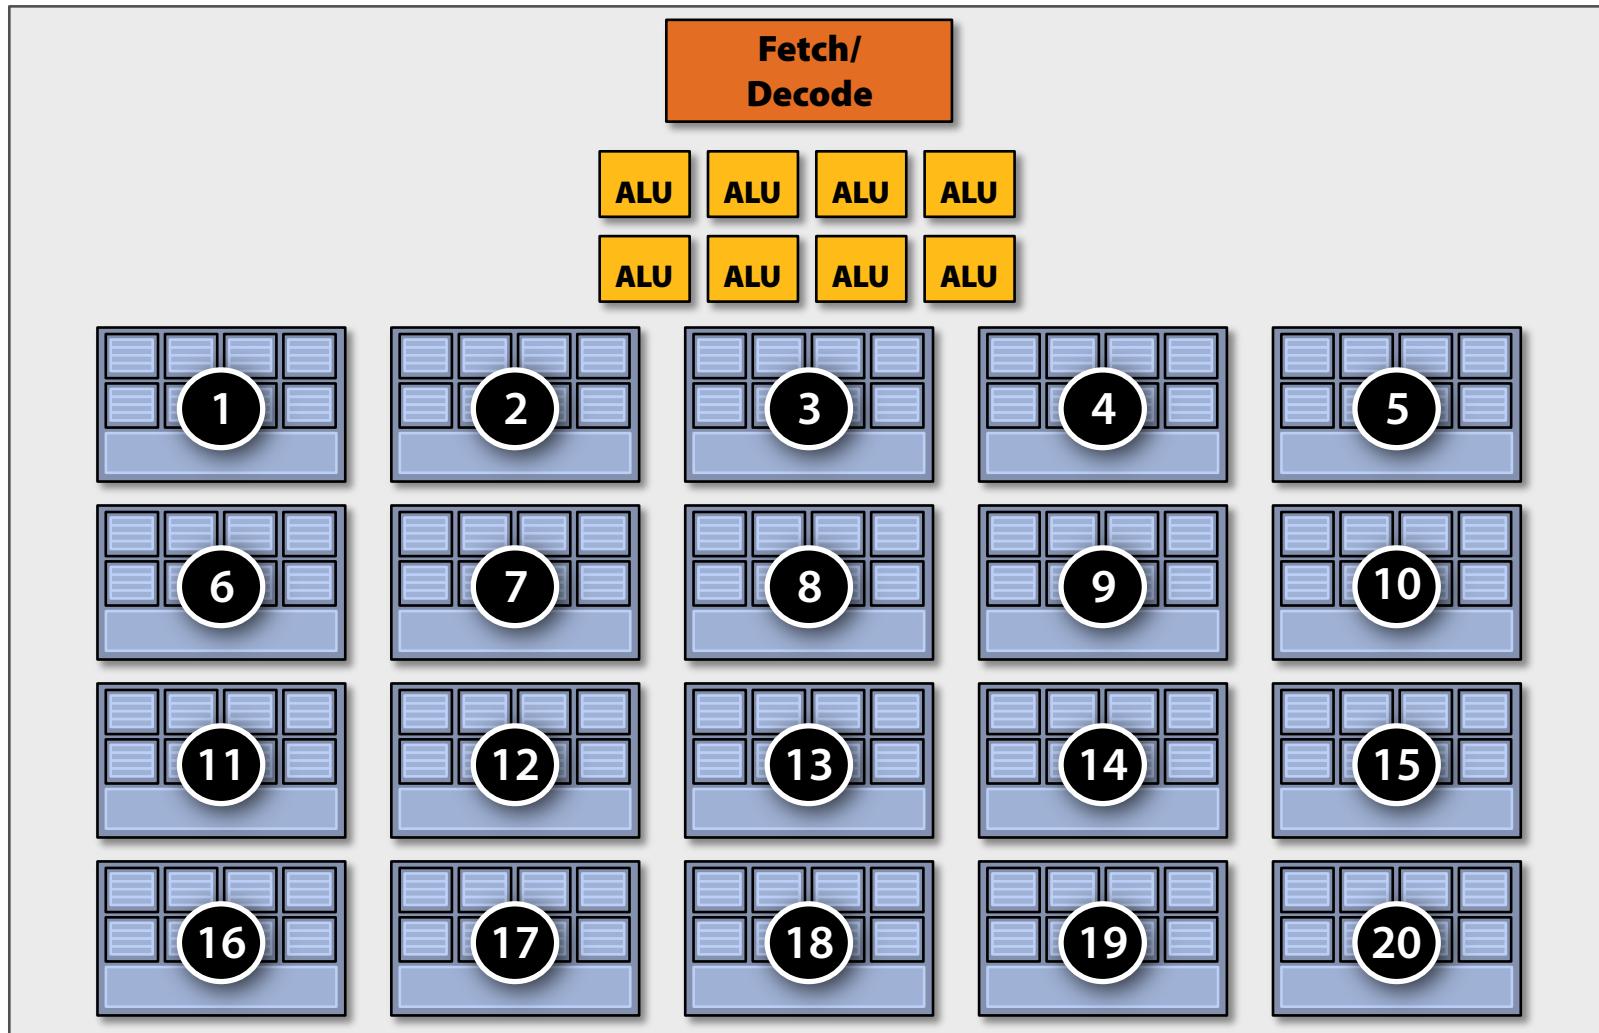




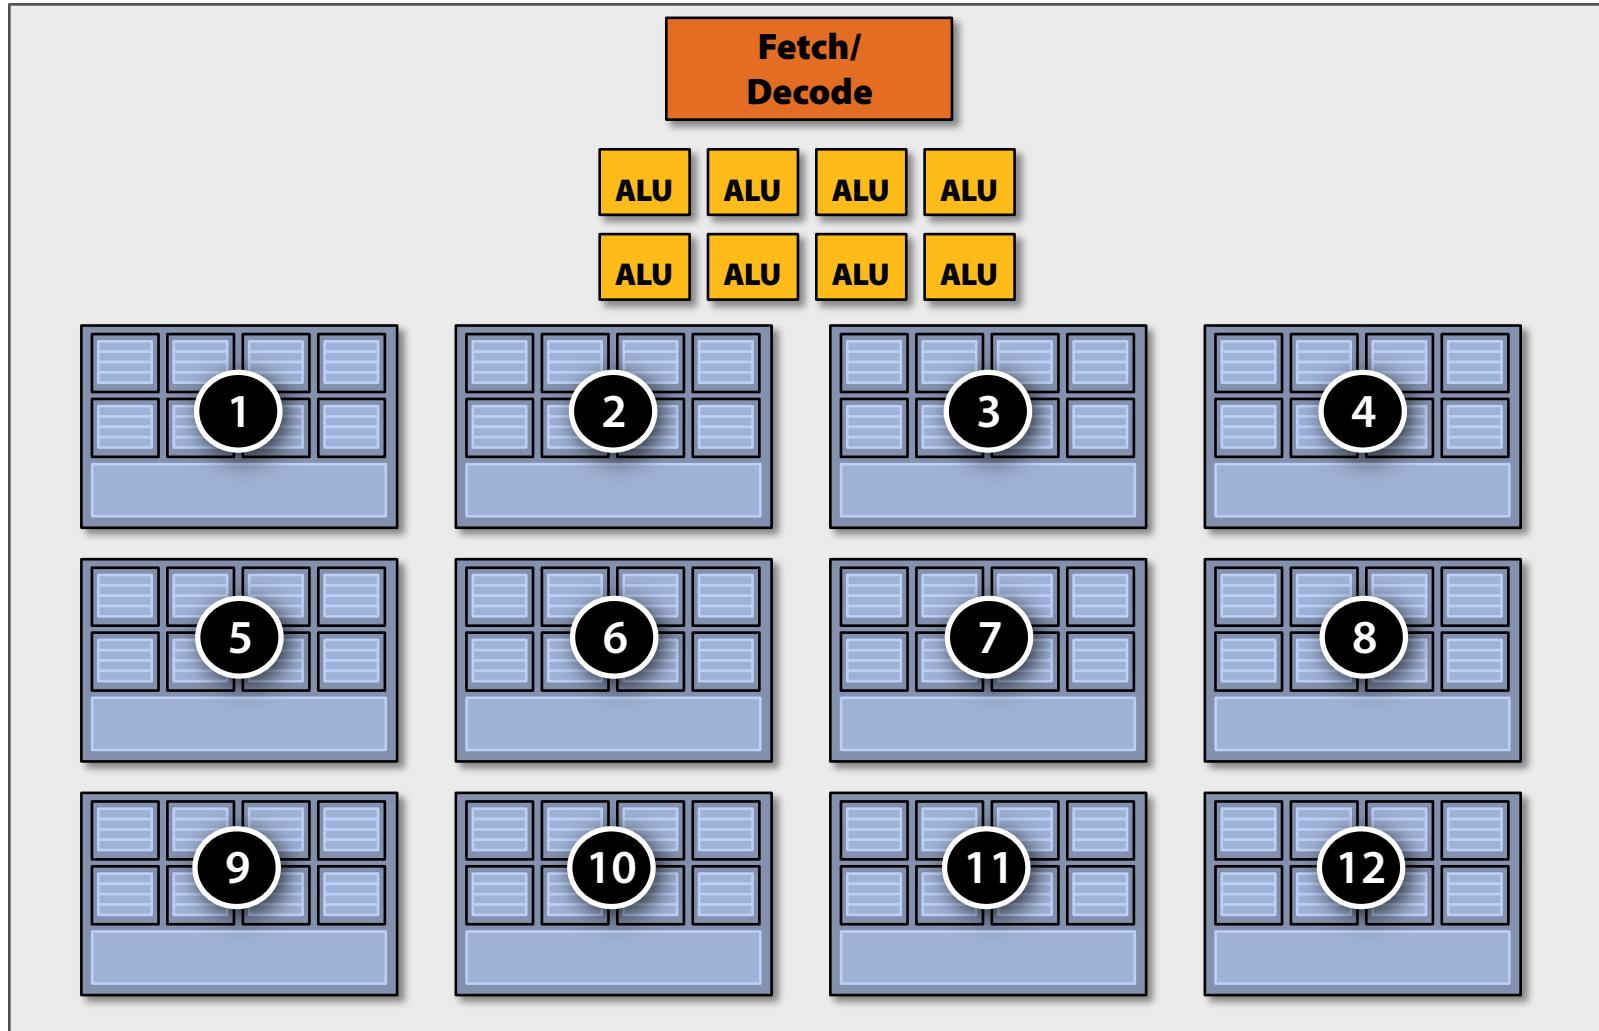
SIGGRAPH2008

Twenty small contexts

(maximal latency hiding ability)

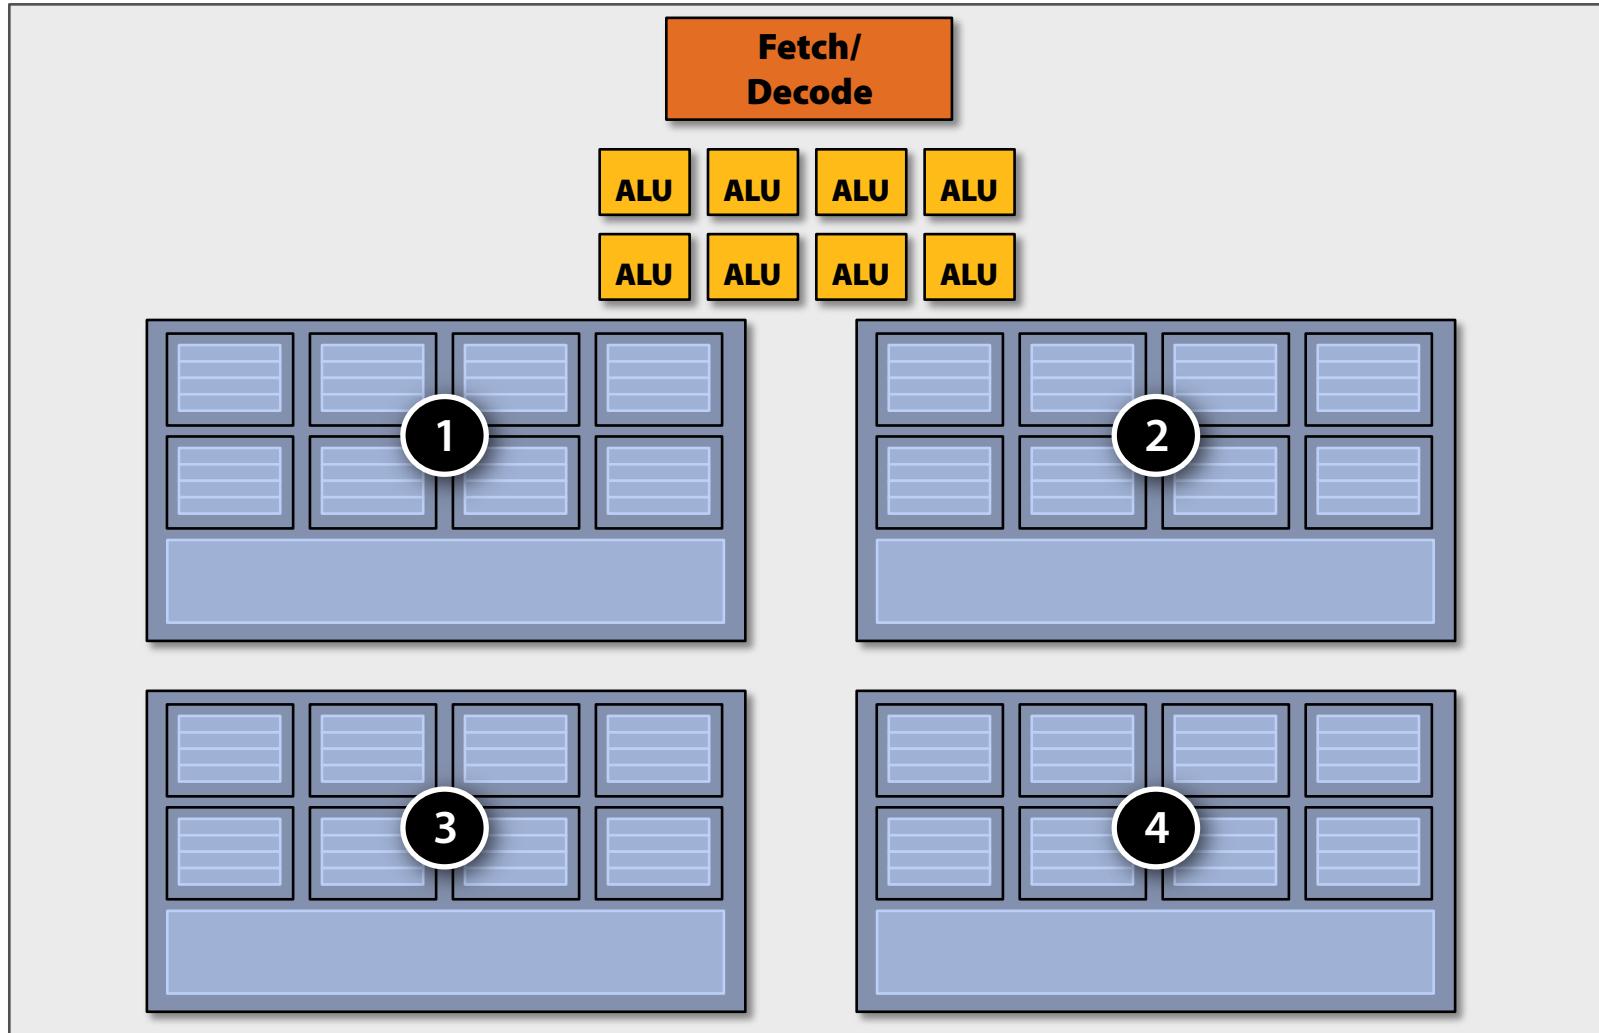


Twelve medium contexts



Four large contexts

(low latency hiding ability)



My chip!

16 cores

8 mul-add ALUs per core
(128 total)

16 simultaneous
instruction streams

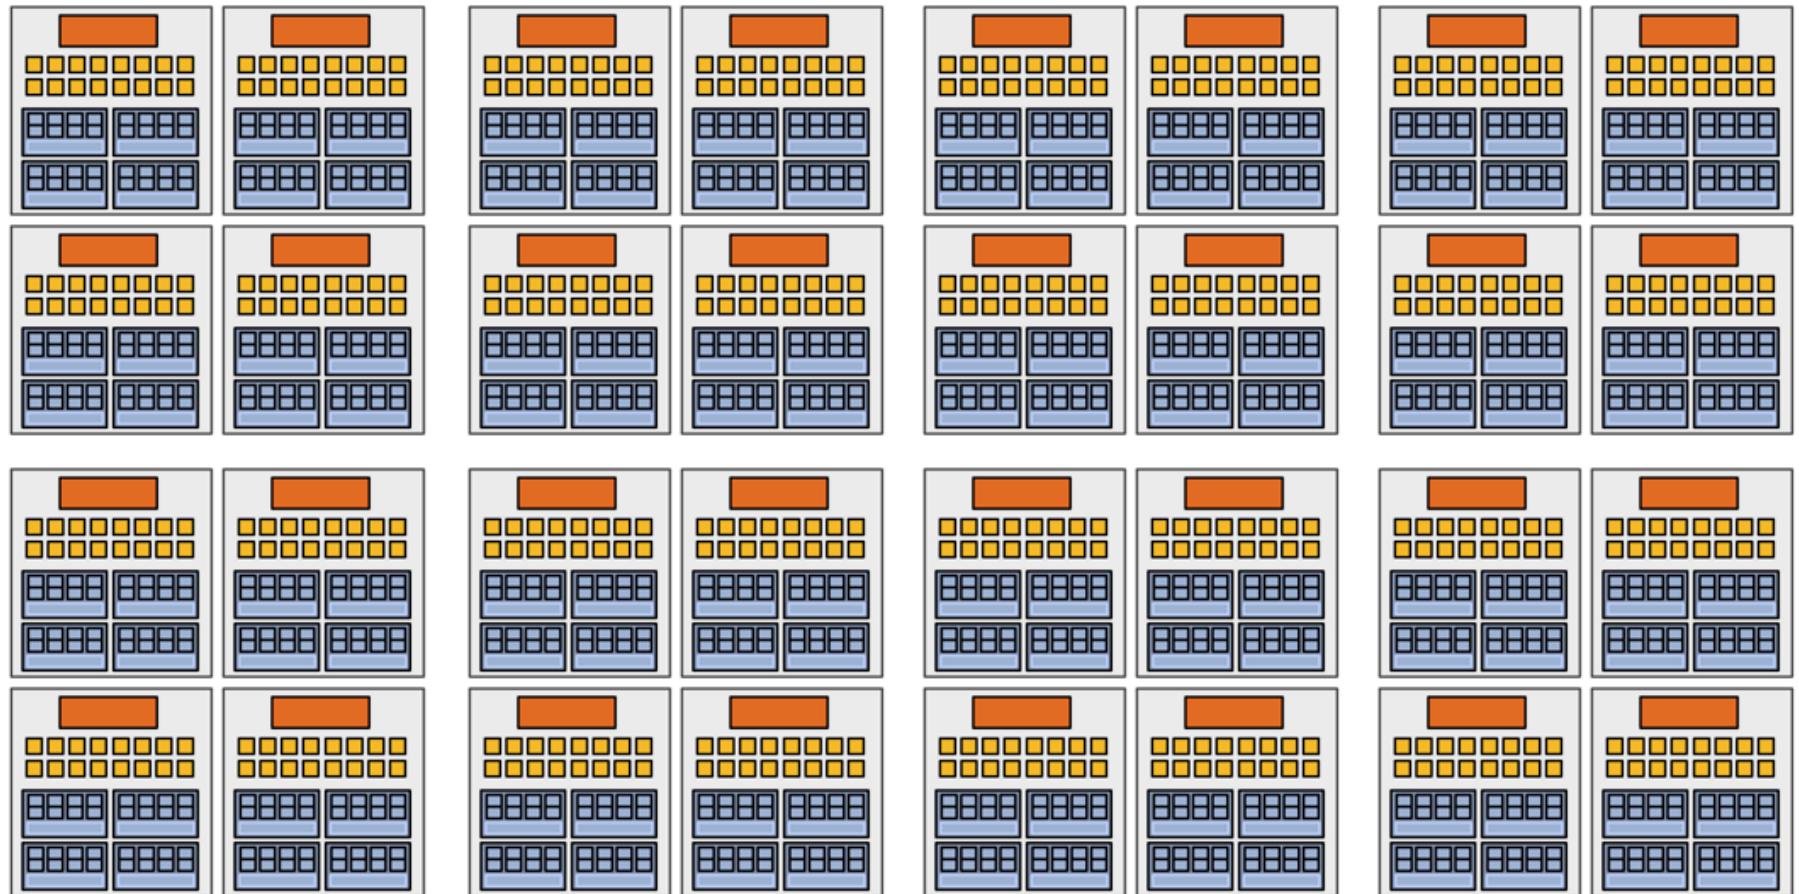
64 concurrent (but interleaved)
instruction streams

512 concurrent fragments

= 256 GFLOPs (@ 1GHz)



My “enthusiast” chip!



32 cores, 16 ALUs per core (512 total) = 1 TFLOP (@ 1 GHz)



Summary: three key ideas

1. Use many “slimmed down cores” to run in parallel
2. Pack cores full of ALUs (by sharing instruction stream across groups of fragments)
 - Option 1: Explicit SIMD vector instructions
 - Option 2: Implicit sharing managed by hardware
3. Avoid latency stalls by interleaving execution of many groups of fragments
 - When one group stalls, work on another group



SIGGRAPH2008

Thank you

Additional information on “supplemental slides” and at
<http://graphics.stanford.edu/~kayvonf/gblog>

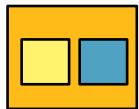
Supplemental slides

- Disclaimer #1: the following slides describe “how I think” about the NVIDIA GTX 280, ATI Radeon 4870, and Intel Larrabee GPUs
- Disclaimer #2: Many other factors play a role in actual chip performance
- These slides use the same hand-wavy notion of “core” as I established in the talk
- Remember: you can substitute the term vertex, primitive, CUDA thread, compute shader thread, or OpenCL work item, for “fragment” (pick your favorite language)

GPU block diagram key



= single “physical” instruction stream fetch/decode
(functional unit control)



= SIMD programmable functional unit (FU), control shared with other functional units. This functional unit may contain multiple 32-bit “ALUs”

- = 32-bit mul-add unit
- = 32-bit multiply unit



= execution context storage



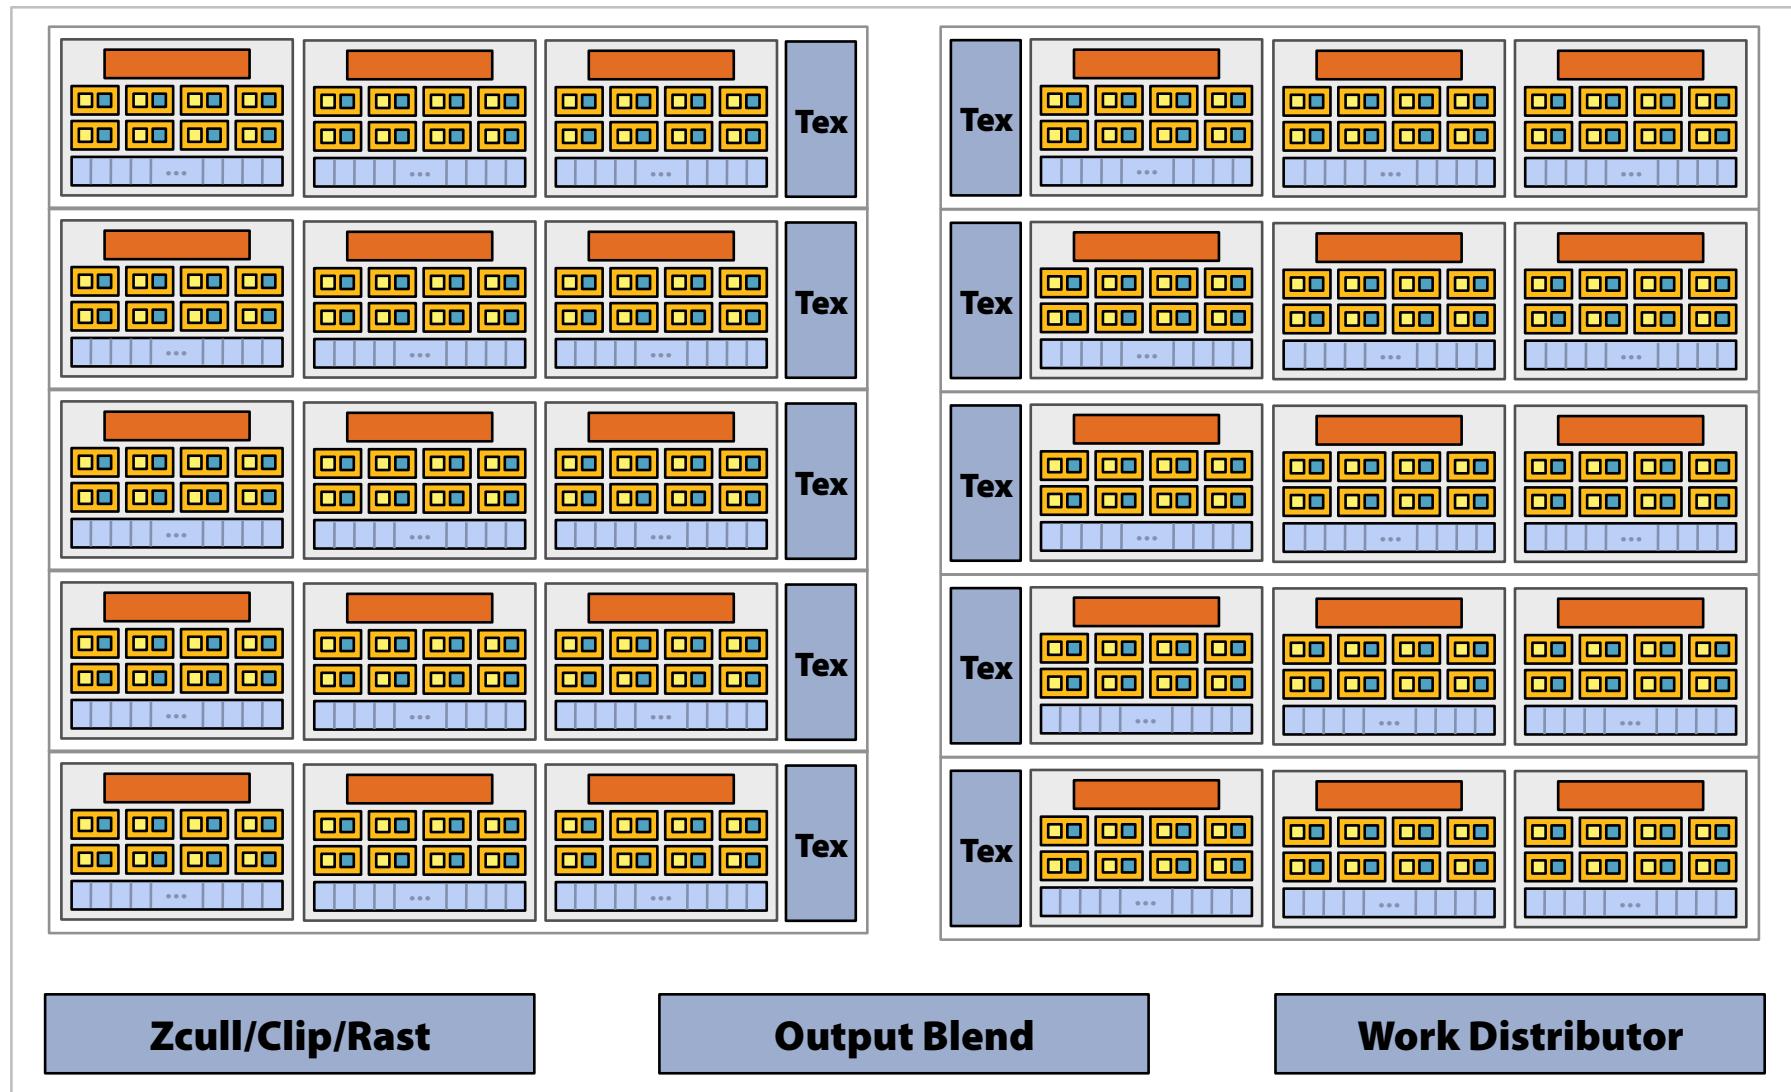
= fixed function unit

NVIDIA GeForce GTX 280



- NVIDIA-speak:
 - 240 stream processors
 - “SIMT execution” (automatic HW-managed sharing of instruction stream)
- Generic speak:
 - 30 processing cores
 - 8 SIMD functional units per core
 - 1 mul-add (2 flops) + 1 mul per functional units (3 flops/clock)
 - Best case: 240 mul-adds + 240 muls per clock
 - 1.3 GHz clock
 - $30 * 8 * (2 + 1) * 1.3 = 933 \text{ GFLOPS}$
- Mapping data-parallelism to chip:
 - Instruction stream shared across 32 fragments (16 for vertices)
 - 8 fragments run on 8 SIMD functional units in one clock
 - Instruction repeated for 4 clocks (2 clocks for vertices)

NVIDIA GeForce GTX 280





SIGGRAPH2008

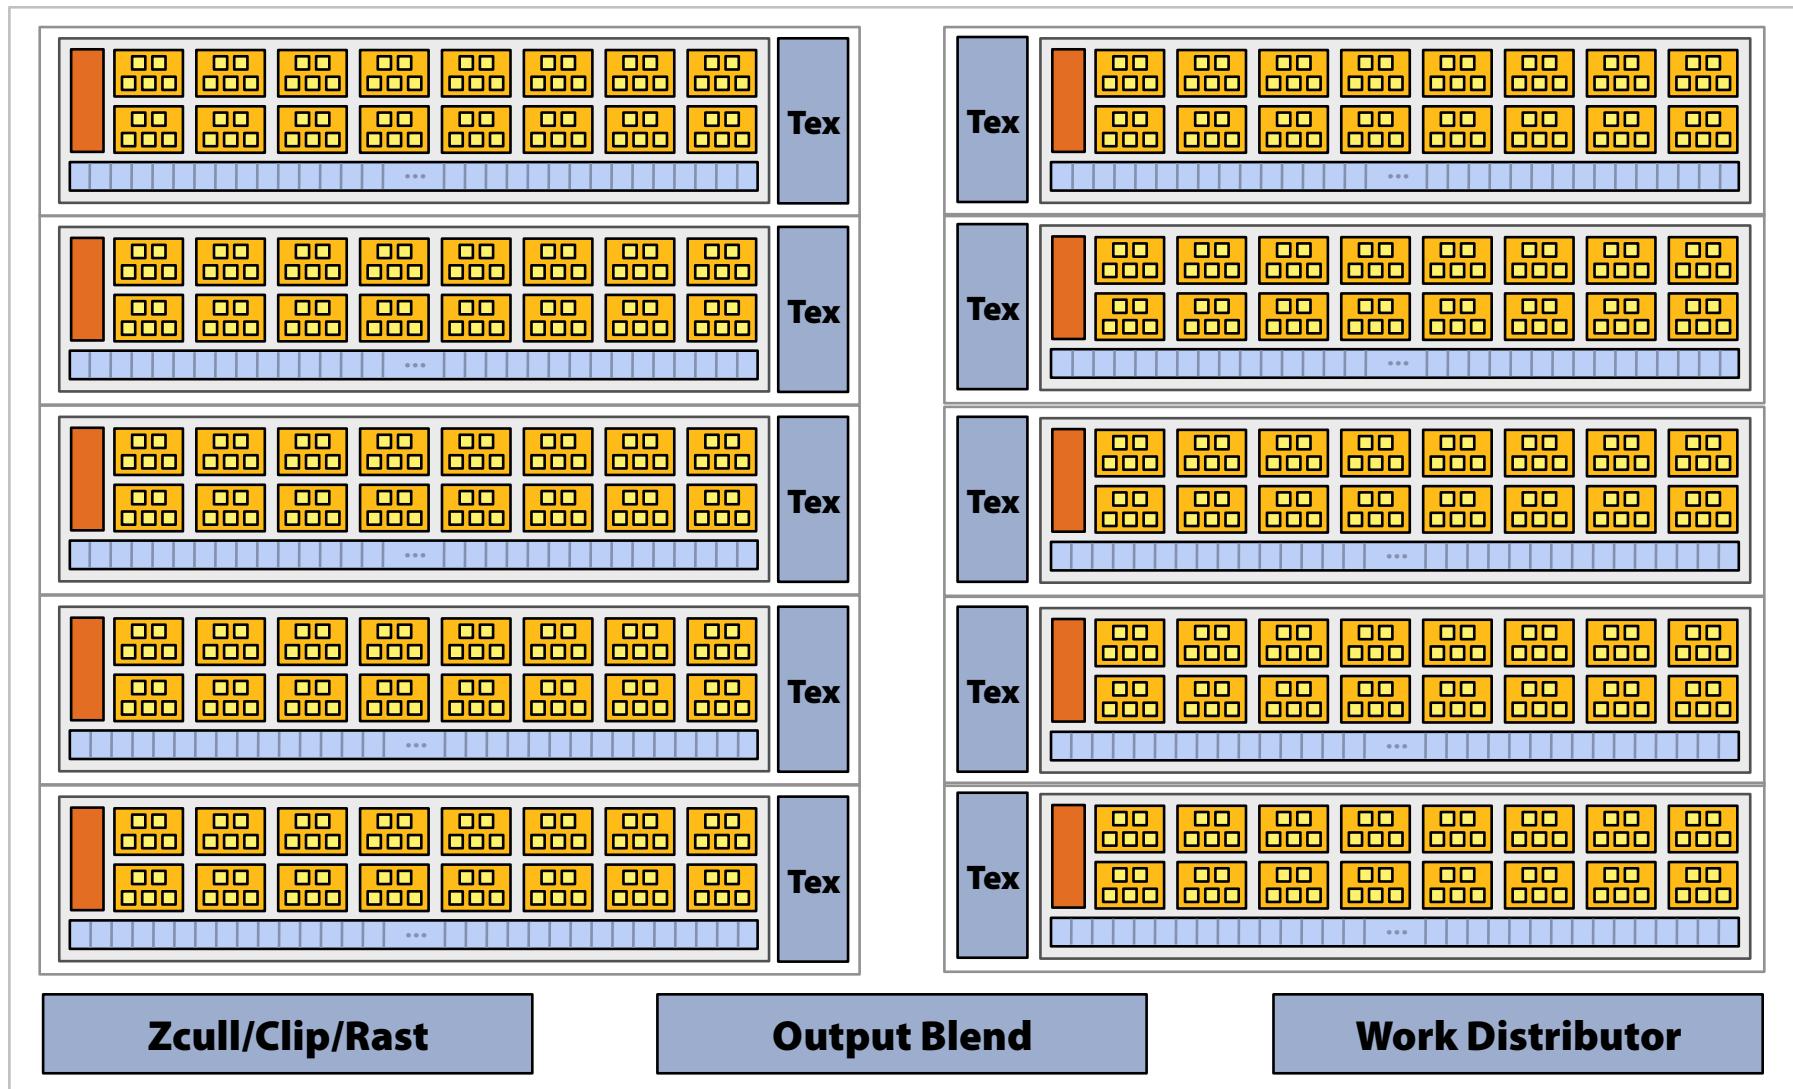
ATI Radeon 4870

- AMD/ATI-speak:
 - 800 stream processors
 - Automatic HW-managed sharing of scalar instruction stream (like “SIMT”)
- Generic speak:
 - 10 processing cores
 - 16 SIMD functional units per core
 - 5 mul-adds per functional unit ($5 * 2 = 10$ flops/clock)
 - Best case: 800 mul-adds per clock
 - 750 MHz clock
 - $10 * 16 * 5 * 2 * .75 = 1.2$ TFLOPS
- Mapping data-parallelism to chip:
 - Instruction stream shared across 64 fragments
 - 16 fragments run on 16 SIMD functional units in one clock
 - Instruction repeated for 4 consecutive clocks

ATI Radeon 4870



SIGGRAPH2008



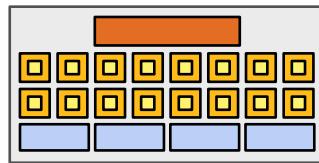
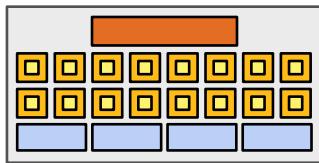


SIGGRAPH2008

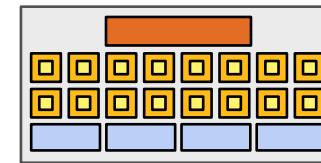
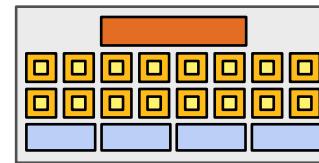
Intel Larrabee

- Intel speak:
 - We won't say anything about the number of cores or clock rate of cores
 - Explicit 16-wide vector ISA
 - If 1GHz clock (then 1 core = 1 LRB unit = 32 GFLOPS from paper)
- Generic speak:
 - X processing cores
 - 16 SIMD functional units per core
 - 1 mul-add per functional unit (2 flops/clock)
 - Best case: 16X mul-adds per clock
 - If you wanted to compete with current GPUs (~ 1 TFLOP), you need about 32 Larrabee units
 - $32 * 16 * 2 = 1 \text{ TFLOP}$
- Mapping data-parallelism to chip:
 - Compilation options determine instruction stream sharing across fragments (a multiple of 16)
 - 16 fragments run on 16 SIMD functional units in one clock

Intel Larrabee (SIGGRAPH paper)



...

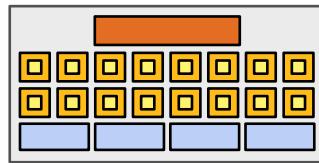
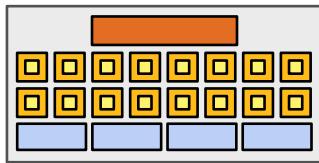


:

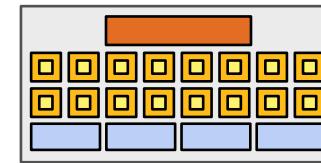
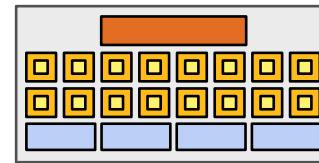
?

?

:



...



... ? ...

